

**Development and demonstration of a proof-of-concept for the integration of programming frameworks for high performance computing into a container-based workflow orchestrator.**

## 2. Project Paper

submitted on November 18, 2023

Faculty for Business and Health

Degree program in Business Informatics

International Management for Business and Information Technology

by

JON ECKERTH

Practice supervisor :

⟨ Hewlett Packard Labs ⟩  
⟨ Harumi Kuno, PhD ⟩  
⟨ Principal Research Scientist ⟩

DHBW Stuttgart:

⟨ Dominic Viola ⟩

Signature of the practice supervisor:

# Statement of Integrity

Hereby I solemnly declare:

1. that this 2. Project Paper titled "*Development and demonstration of a proof-of-concept for the integration of programming frameworks for high performance computing into a container-based workflow orchestrator.*" is entirely the product of my own scholarly work, unless otherwise indicated in the text or references, or acknowledged below;
2. I have indicated the thoughts adopted directly or indirectly from other sources at the appropriate places within the document;
3. this 2. Project Paper has not been submitted either in whole or part, for a degree at this or any other university or institution;
4. I have not published this 2. Project Paper in the past;
5. the printed version is equivalent to the submitted electronic one.

I am aware that a dishonest declaration will entail legal consequences.

Dublin, November 18, 2023



**Jon Eckerth**

**Abstract:**

This Second project paper presents the development and demonstration of a proof-of-concept for the integration of High Performance Computing (HPC) programming frameworks into a container-base workflow orchestrator.

This project aim to reconcile the different demands of the HPC and Cloud Computing (CC) communities by demonstrating the integration of the HPC programming framework Arkouda into the container-based workflow orchestrator Pachyderm, showing the technical feasibility of this approach.

A prototype implementing this integrated system is constructed and evaluated through prototyping methodologies, with a focus on resilience, scalability, portability, and user-friendliness. The prototype is iteratively refined to address Loosely Coupled Problems (LCP) and Tightly Coupled Problems (TCP), with particular attention to the usability of the system for non CC experts.

This project paper contributes to the body of knowledge by way of practical example, lessons learned with each iteration, and sheds light on pathways for future research towards a landscape where the seamless and efficient integration of HPC workloads in CC environments becomes possible

# Contents

|                                                         |            |
|---------------------------------------------------------|------------|
| <b>List of abbreviations</b>                            | <b>IV</b>  |
| <b>List of figures</b>                                  | <b>VI</b>  |
| <b>List of tables</b>                                   | <b>VII</b> |
| <b>1 Introduction</b>                                   | <b>1</b>   |
| 1.1 Motivation . . . . .                                | 1          |
| 1.2 Problem Statement . . . . .                         | 2          |
| 1.3 Research Questions . . . . .                        | 3          |
| 1.4 Contributions . . . . .                             | 4          |
| <b>2 Methodology</b>                                    | <b>4</b>   |
| 2.1 Prototyping . . . . .                               | 4          |
| 2.2 Decision-Making . . . . .                           | 5          |
| 2.2.1 Weighted Sum Model . . . . .                      | 5          |
| 2.2.2 Simple Multi-Attribute Rating Technique . . . . . | 6          |
| 2.2.3 SMART Exploiting Ranks . . . . .                  | 7          |
| <b>3 State of the Art</b>                               | <b>8</b>   |
| 3.1 Containerization . . . . .                          | 8          |
| 3.2 High Performance Computing Frameworks . . . . .     | 10         |
| 3.2.1 Loosely Coupled Problems . . . . .                | 10         |
| 3.2.2 Tightly Coupled Problems . . . . .                | 11         |
| <b>4 Creation of the Artifact</b>                       | <b>11</b>  |
| 4.1 Initial Goals . . . . .                             | 11         |
| 4.2 Overall Structure . . . . .                         | 12         |
| 4.3 Selection of Workflow Management Tools . . . . .    | 12         |
| 4.4 Implementation of the Artifact . . . . .            | 17         |
| 4.4.1 Infrastructure . . . . .                          | 17         |
| 4.4.2 Tightly Coupled HPC Workloads . . . . .           | 20         |
| 4.4.3 Supplementary Services . . . . .                  | 24         |
| 4.5 Evaluation of the Artifact . . . . .                | 28         |
| <b>5 Summary and Outlook</b>                            | <b>29</b>  |
| 5.1 Summary . . . . .                                   | 29         |
| 5.2 Outlook . . . . .                                   | 30         |
| <b>Appendix</b>                                         | <b>31</b>  |
| <b>List of literature</b>                               | <b>54</b>  |

# List of abbreviations

|                  |                                            |
|------------------|--------------------------------------------|
| <b>AI</b>        | Artificial Intelligence                    |
| <b>CC</b>        | Cloud Computing                            |
| <b>CI/CD</b>     | Continuous Integration/Continuous Delivery |
| <b>CLASP</b>     | Cloud Application Services Platform        |
| <b>CNCF</b>      | Cloud Native Computing Foundation          |
| <b>CNI</b>       | Container Network Interface                |
| <b>CWL</b>       | Common Workflow Language                   |
| <b>FAM</b>       | Fabric Attached Memory                     |
| <b>FOSS</b>      | Free and Open Source Software              |
| <b>GASNet</b>    | Global Address Space Networking            |
| <b>HPC</b>       | High Performance Computing                 |
| <b>HPE</b>       | Hewlett Packard Enterprise                 |
| <b>IaC</b>       | Infrastructure as Code                     |
| <b>IP</b>        | Intellectual Property                      |
| <b>k8s</b>       | Kubernetes                                 |
| <b>LCP</b>       | Loosely Coupled Problems                   |
| <b>MAUT</b>      | Multi-attribute Utility Theory             |
| <b>ML</b>        | Machine Learning                           |
| <b>MPI</b>       | Message Passing Interface                  |
| <b>NPO</b>       | Non-Governmental Organization              |
| <b>OpenSHMEM</b> | Open Shared Memory                         |
| <b>PFS</b>       | Pachyderm File System                      |
| <b>PoC</b>       | Proof of Concept                           |
| <b>PV</b>        | Persistent Volume                          |
| <b>RAD</b>       | Rapid Application Development              |
| <b>RBAC</b>      | Role Based Access Control                  |
| <b>RDMA</b>      | Remote Direct Memory Access                |
| <b>ROC</b>       | Rank order centroid                        |
| <b>RR</b>        | Rank Reciprocal                            |

|                 |                                        |
|-----------------|----------------------------------------|
| <b>RS</b>       | Rank Sum                               |
| <b>SMART-ER</b> | SMART Exploiting Ranks                 |
| <b>SMART</b>    | Simple Multiattribute Rating Technique |
| <b>SME</b>      | Subject-Matter Expert                  |
| <b>SSO</b>      | Singele Sign On                        |
| <b>TCP</b>      | Tightly Coupled Problems               |
| <b>UDP</b>      | User Datagram Protocol                 |
| <b>VM</b>       | Virtual Machine                        |
| <b>WSM</b>      | Weighted Sum Model                     |

## List of Figures

|    |                                                                                                         |    |
|----|---------------------------------------------------------------------------------------------------------|----|
| 1  | Formula for calculating the Weighted Sum Model (WSM) score <sup>1</sup> . . . . .                       | 6  |
| 2  | Formula for calculating the Simple Multiattribute Rating Technique (SMART) score <sup>2</sup> . . . . . | 7  |
| 3  | Formula for the Rank order centroid (ROC) weights . . . . .                                             | 7  |
| 4  | Formula for the Rank Sum (RS) weights . . . . .                                                         | 8  |
| 5  | Formula for the Rank Reciprocal (RR) weights . . . . .                                                  | 8  |
| 6  | Pachykouda high level diagram showing three main aspects . . . . .                                      | 12 |
| 7  | Pachykouda datum distribution amongst workers . . . . .                                                 | 20 |
| 8  | Arkouda workers on Kubernetes (k8s) . . . . .                                                           | 22 |
| 9  | Pachyderm High-Level Architecture . . . . .                                                             | 24 |
| 10 | Stars and Stack overflow Questions Comparison . . . . .                                                 | 32 |
| 11 | Maturity Metric comparison (higher rated projects) . . . . .                                            | 37 |
| 12 | Maturity Metric comparison (lower rated projects) . . . . .                                             | 37 |
| 13 | ROC weights . . . . .                                                                                   | 38 |
| 14 | RR weights . . . . .                                                                                    | 39 |
| 15 | RS weights . . . . .                                                                                    | 39 |
| 16 | Swim lane Diagram of the communication between the user and Pachyderm . . . . .                         | 40 |
| 17 | Flowchart of the Jenkins Pipeline . . . . .                                                             | 41 |
| 18 | Continued flowchart of the Jenkins Pipeline . . . . .                                                   | 42 |

## List of Tables

|    |                                                                        |    |
|----|------------------------------------------------------------------------|----|
| 1  | Weighting of the criteria . . . . .                                    | 15 |
| 2  | Evaluation of the suggested tools . . . . .                            | 15 |
| 3  | Evaluation of the additional tools . . . . .                           | 16 |
| 4  | Overall score of the tools . . . . .                                   | 16 |
| 5  | Community Support Data . . . . .                                       | 32 |
| 6  | Composite scores of Workflow managers, sorted by final score . . . . . | 33 |
| 7  | Strategic Alignment Scores . . . . .                                   | 36 |
| 8  | Maturity Data . . . . .                                                | 36 |
| 9  | Maturity Scores . . . . .                                              | 37 |
| 10 | Scores Based on Cost . . . . .                                         | 38 |



# 1 Introduction

In this section, the underlying motivation of this project is explained. Furthermore, the problems which will be addressed by this project are described, which serve as the basis for the research questions which will guide this project and ultimately result in solutions and further questions which are listed in the contributions section and discussed in the conclusion.

## 1.1 Motivation

The proliferation of "Big Data" has led to the need to compute, analyze, and visualize ever-increasing amounts of datasets, which themselves are getting more and more complex. This has led to an ever-increasing demand for more efficient and quicker ways to process data.

Both the HPC and the CC community have been working on solutions to distribute and parallelize computations for decades, both with their own approaches and solutions to their respective problems.

While the HPC community has been putting a lot of effort into developing new and extremely efficient ways to parallelize computations, the CC community has been focusing on improving the flexibility, scalability and resilience of their solutions as well as enhancing usability for developers and end-users.

Both used to be very distinct and separate communities due to their very different use cases. While the HPC community was mostly concerned with scientific computing and simulations of physical phenomena, the CC community is mostly concerned with providing a reliable and easily up and down scalable infrastructure for the industry and businesses.

Now with the advent of Machine Learning (ML) and Artificial Intelligence (AI) the two communities are starting to converge, as the ML and AI community is adopting the tools and techniques of both communities to solve their problems as they see fit.

However this convergence of the two is not without its problems, being developed in two separate, coexisting communities, the tools and techniques of both communities are not always compatible with each other. The goal of this project is to find a way to bridge this gap and to find a way to combine the best of both worlds.

## 1.2 Problem Statement

The following key problems have emerged from the convergence of High Performance Computing (HPC) and Cloud Computing (CC) communities, especially in the context of Machine Learning (ML) and Artificial Intelligence (AI) research:

- **Workload Resilience and Fault Tolerance in HPC:** HPC systems often lack mechanisms to recover from task failures within larger jobs, running for an extended time. When a component task fails, it can invalidate the entire computation, requiring a restart from scratch. This need for resilient failover and verification strategies as well as the need to avoid computational wastage is a key challenge for HPC systems, especially with ever-increasing system sizes and complexity.<sup>3</sup>
- **Environment/Package Management in HPC:** HPC systems are notorious for their complex package management systems. As having a shared infrastructure between many users, each with their own specific needs and requirements of different versions of packages, libraries and software, all the while sharing a common environment. Many solutions to this problem have been developed, each with their own advantages and disadvantages.<sup>4567</sup>
- **Portability Issues with HPC:** HPC systems, while designed for specific hardware and software stacks, can technically support application migration. Yet, achieving comparable performance on different HPC systems can be problematic, as optimized libraries are crucial for high efficiency<sup>8</sup>. Without these, performance may significantly decline. Contrastingly, CC environments favor a platform-agnostic approach with containerization to aid portability. Nonetheless, whether CC preserves performance consistently remains a question. Moreover, optimizing applications on HPC often requires administrative input to access necessary libraries and environment settings.
- **Scalability and Flexibility in HPC:** Due to its direct access to the hardware and very specific hardware needs, HPC systems are often inflexible and hard to dynamically scale. While CC systems are designed to be easily scalable and flexible and are often designed to be hardware-agnostic and abstract away the underlying hardware. This becomes especially relevant in the context of heterogeneous hardware, where the hardware is not uniform and consists of different types of hardware, which is becoming more and more common in the context of ML and AI research.
- **Lack of Interconnected Problem-Solving in CC:** The workloads traditionally deployed on CC systems are often independent of each other, like load balancing, web hosting, etc. This is in stark contrast to the interconnected nature of HPC workloads, where

---

<sup>3</sup>Egwutuoha et al. 2013, pp. 1302–1326

<sup>4</sup>Dubois/Epperly/Kumfert 2003, pp. 83–88

<sup>5</sup>Bzeznik et al. 2017, pp. 1–6

<sup>6</sup>Gamblin et al. 2015, pp. 1–12

<sup>7</sup>Hoste et al. 2012, pp. 572–582

<sup>8</sup>Canon/Younge 2019, pp. 49–54

each part of the input data may depend on the other parts of the input data, such that all nodes of the system need to be able to communicate with each other.

- **Provenance and Reproducibility:** Another need that is becoming more and more important in the context of ML and AI research is the need for provenance and reproducibility of results. Being able to tell which data was used to train the model, is of ever-increasing importance as the influence the resulting models have on our lives increases as well as the data used to train the model. This is especially important since it is crucial to ensure that the data is not biased, outdated, or otherwise flawed, which could lead to incorrect predictions, decisions, or recommendations. In addition various data sources, from images to text, may have copyright restrictions that, when overlooked, can lead to significant legal complications.
- **Versioning Limitations:** The dynamic nature of ML and AI research necessitates robust versioning solutions for data, configurations and code. CC has developed many solutions to this problem over the years, making them their own subsection of the ecosystem, namely Continuous Integration/Continuous Delivery (CI/CD) tools for the testing and deployment of applications as well as Infrastructure as Code (IaC) tools for the deployment of infrastructure. While many solutions have been developed for the one-off deployment of HPC systems, the dynamic nature of CC systems necessitates a more robust solution to this problem, from which the HPC community could benefit as well.

### 1.3 Research Questions

To address the aforementioned problems, to bridge the gap between the two paradigms and to combine the best of both worlds, an integration of the two paradigms is needed. This was accomplished by integrating a HPC framework called Arkouda<sup>9</sup> into a container based CC workflow management tool called 'Pachyderm'<sup>10</sup> and integrating both with the supporting infrastructure the CC system enables us to use. This process of integration and prototyping as well as the explanation of the underlying concepts and technologies will be the focus of this project.

- **RQ1:** *How can a high-performance computing framework be effectively integrated into a container-based workflow management tool?*
- **RQ2:** *What are the opportunities for improving the integration of high-performance computing frameworks with container-based workflow management tools?*

---

<sup>9</sup>Merrill/Reus/Neumann 2019, p. 28

<sup>10</sup>Pachyderm 2022a

## 1.4 Contributions

In order to address the problems stated above, find answers to the research questions and to bridge the gap between the two paradigms, the following contributions were made:

- **C1:** *An analysis of the problem space and existing solution, within the constraints of time, resources, and businesses needs.*
- **C2:** *A prototype implementation combining the Arkouda framework with the k8s based workflow orchestrator 'Pachyderm' running on the Heydar Cluster.*
- **C3:** *Further integrations of tools from both sides of the spectrum, addressing many of aforementioned pain-points.*

## 2 Methodology

### 2.1 Prototyping

The methodology of Prototyping in the context of software development has been a well established practice for many years and has been discussed in literature for at least the last 40 years<sup>11</sup>. It is especially useful in complex situations where the requirements might not be fully known or understood at the beginning of the project and need to be discovered in an iterative and exploratory process, as it enables the project to be partitioned into more manageable chunks which can be tackled individually<sup>12</sup>.

At their core prototyping is a state transition model, where a final state is approached in a series of non-ideal intermediate steps, each of which revealing new information about the problem domain and the solution space and applying this new information to the next iteration<sup>13</sup>.

It is described as the methodology of the highest constructive plasticity, as it supports many different flavors and variations of the core concept. As long as it follows the core principles of this iterative approximation of an ideal version, and specifically evaluates the results of the iterations it should be considered a methodologically complete scientific method<sup>14</sup>.

For this project we will be using a concatenated version of the so called Rapid Application Development (RAD) model, which has been quite established in software development, as it is very well suited for projects with a high degree of complexity and uncertainty, while reducing the administrative overhead of the project<sup>15</sup>. Reducing the administrative overhead is especially

---

<sup>11</sup>Gomaa 1983, pp. 17–27

<sup>12</sup>Naumann/Jenkins, A. M. 1982, pp. 29–44

<sup>13</sup>Kraushaar/Shirland 1985, pp. 189–197

<sup>14</sup>Wilde/Hess 2007, pp. 280–287

<sup>15</sup>Martin 1991

important for this project, as the time frame is limited to 3 months, and the project is being done by a single person, therefore more complex tools like the spiral model<sup>16</sup>, which is a popular model emphasizing a risk analysis with each step of the iteration, would be too time-consuming to be feasible for this project.

First we define a loose goal for the project, which will be refined and adjusted with each iteration as we learn more about the problem domain and the solution space. Because of this we do not need to plan as meticulously as we would for a waterfall approach, where the requirements are known at the beginning of the project.

We then enter the section with the prototyping iterations, where we will alternate between the implementation of what we planned in the previous section and the evaluation of the results of the previous iteration which serve as the basis for the next iteration.

Then a final evaluation of the whole project will be done, where the final results will be evaluated and the project will be concluded, and advice for the pickup of the project will be given.

## 2.2 Decision-Making

As previously described, the methodology of Prototyping benefits from a very tight loop of iterations between the different phases of the project. While this is highly effective in producing a good end result, it can also take many iterations and a lot of experimentation until an adequate tool or solution has been found. Given the constraints of a limited time frame for this project, it becomes crucial to use this time as efficiently as possible. Sometimes, when the time does not permit a thorough exploration of all the possible options, a decision with incomplete information needs to be made.

To ensure that the decisions made are the most optimal within the constraints of the available information, adopting a systematic, replicable, and transparent decision-making process becomes essential. Over the years, various frameworks have been crafted to guide decision-making, particularly when information is complex and multidimensional.

### 2.2.1 Weighted Sum Model

Evangelos Triantaphyllou suggests that the Weighted Sum Model (WSM) is in practice the most used and most relevant decision-making framework<sup>17</sup>. The WSM method, by design, mandates the assignment of specific weights to each criterion based on its relevance. Subsequent to this, every alternative is evaluated based on these weighted criteria, resulting in a cumulative score. The alternative with the highest score is therefore the optimal choice.

---

<sup>16</sup>Boehm 1988, pp. 61–72

<sup>17</sup>Triantaphyllou 2000, pp. 1–4, p. 1

<sup>18</sup>*Weighted Sum Model* 2022

$$A_i^{WSM-score} = \sum_{j=1}^n w_j a_{ij} \quad \text{for } i = 1, 2, 3, \dots, m.$$

Abb. 1: Formula for calculating the WSM score<sup>18</sup>

Where:

- $w_j$ : This represents the weight assigned to the  $j$ -th criterion. Weights are determined by the decision-makers based on the relative importance of each criterion. They should be normalized (i.e., the sum of all weights should be 1 or 100%) to maintain a consistent scale.
- $a_{ij}$ : This represents the score or rating of the  $i$ -th alternative concerning the  $j$ -th criterion. This score is an assessment of how well the alternative meets or satisfies the specific criterion.

This method, despite its simplicity and direct approach, isn't without limitations. One notable drawback is its dependence on dimensionless scales. For the weights to properly reflect the criteria's importance, the scores need to be on a common, dimensionless scale, a detail not always feasible or convenient in practice.

### 2.2.2 Simple Multi-Attribute Rating Technique

In contrast to the WSM, which predominantly utilizes a direct mathematical approach to rank alternatives based on their weighted sum scores, the Simple Multiattribute Rating Technique (SMART) methodology offers a more comprehensive approach to multi-criteria decision-making. While WSM is primarily concerned with simple weighted arithmetic sums, the SMART method dives deeper, ensuring that diverse performance values—both quantitative and qualitative are harmonized and placed on a common scale.

The SMART method, grounded in Multi-attribute Utility Theory (MAUT), provides a structured framework that encompasses more than just the weighting of criteria. It involves:

1. Discernment of vital criteria pertinent to the decision in focus.
2. Weight allocation to each criterion in accordance to its significance.
3. Evaluation of each potential alternative against the identified criteria, culminating in a score.
4. Aggregation of these individual scores via their associated weights, yielding a total score for every alternative.

By adhering to the SMART framework, alternatives can be sequenced based on their aggregated weighted scores. This systematic approach equips decision-makers to choose solutions that align closely with their objectives. The computational formula integral to the SMART method is:

$$x_j = \frac{\sum_{i=1}^m w_i a_{ij}}{\sum_{i=1}^m w_i}, \quad j = 1, \dots, n.$$

Abb. 2: Formula for calculating the SMART score<sup>19</sup>

- $x_j$  Is the overall utility score for alternative  $j$ . The higher the score, the better the alternative, in comparison to the other alternatives.
- $a_{ij}$  Is the utility score for alternative  $j$  for the criterion  $i$ .
- $w_i$  Is the weight of criterion  $i$ .

This method's emphasis on utility functions ensures a more nuanced and adaptable approach to decision-making compared to models like WSM, making it suitable for complex scenarios where criteria and alternatives are diverse in nature<sup>20</sup>.

### 2.2.3 SMART Exploiting Ranks

The SMART Exploiting Ranks (SMART-ER) method is a variant of the SMART method that attempts to alleviate the largest issue of the original SMART method, namely the problem of a somewhat arbitrary ranking of the options if no numerical values can be derived.

This method addresses the issue by letting the decision maker simply rank the different criteria in relation to each other and then normalizing the weights<sup>21</sup>. They propose the different weighting curves:

$$w_i(ROC) = \frac{1}{n} \sum_{j=1}^n \frac{1}{j}, \quad i = 1, \dots, n.$$

Abb. 3: Formula for the ROC weights

The ROC takes the centroid of the rank order and uses the reciprocal of the rank as the weight.

The RS uses linear curve where weights are normaliz by dividing them by the sum of all weights.

---

<sup>19</sup>Taken from Fülöp 2005, p. 6

<sup>20</sup>Fülöp 2005, pp. 1–15, p. 6

<sup>20</sup>Barfod/Leleur 2014, p. 26

<sup>21</sup>Roberts/Goodwin 2002, pp. 291–303, p. 296

$$w_i(RS) = \frac{n+1-i}{n(n+1)/2}, \quad i = 1, \dots, n.$$

Abb. 4: Formula for the RS weights

$$w_i(RR) = \left( \frac{\frac{1}{i}}{\sum_{j=1}^n \frac{1}{j}} \right), \quad \text{rank } i = 1, \dots, n, \quad \text{option } j = 1, \dots, n$$

Abb. 5: Formula for the RR weights

The RR emphasizes the most important criteria by using the reciprocal of the rank as the weight, then normalizing each weight by the sum of all reciprocals.

## 3 State of the Art

While this project assumes that the reader is already familiar with the basic concepts of High Performance Computation as well as having a basic understanding of the Cloud Computing technology stack, especially containerization and software defined infrastructure, this chapter will give a brief overview of the current state of the art in the field of High Performance Computing and Cloud Computing.

### 3.1 Containerization

Containerization, a concept originated within the cloud computing community, refers to the process of isolating a process from the rest of the operating system’s user space. It provides a virtual environment for the process to run in, which can be specifically tailored to the process’s needs without impacting the rest of the system. Such isolation is particularly advantageous when running multiple processes on the same machine, as it eliminates the risk of interference between processes or with the system itself<sup>22</sup>. A notable benefit of containerization is its streamlined dependency management, enabling different versions of the same library to coexist on a single machine, with each process accessing the version it requires. Containerization offers a lightweight alternative to traditional virtual machines by building on the host’s kernel to isolate only the user space, rather than emulating an entire operating system<sup>23</sup>. Although prevalent in cloud computing, containerization technologies are only recently being adopted in HPC domains. The HPC community’s cautious approach towards containerization stems from the imperative to preserve

---

<sup>22</sup>Shenoi/Shah/Joshi 2019, pp. 64–70

<sup>23</sup>Docker 2023



system performance, which is paramount in high-stakes computational tasks. This section will delve into the nuances of containerization within the HPC ecosystem, elucidating the reasons behind its gradual integration and the potential benefits it heralds for future applications.

## Container Solutions

Currently the containerization paradigm is spreading rapidly with multiple solutions available, the most popular being Docker<sup>24</sup>. However, within the HPC community, Singularity<sup>25</sup> has gained more and more traction, while the previous options have historically been more popular in CC. The main differentiating factor between the two is that Docker is reliant on the Docker daemon, which needs to be run as root, while Singularity does not need a daemon and can be run as unprivileged user. This makes Singularity generally more appealing to HPC users, as it is more in line with the security policies of most HPC clusters, and simplifies the process of running containers on HPC clusters. But Docker is still the more popular option in CC, as it is more flexible and has a larger ecosystem of tools and services built around it which makes rapid prototyping and development easier.

These tools make use of two features of the Linux kernel, namely cgroups and namespaces<sup>26</sup>. Cgroups are used to limit the resources a process can use, while namespaces are used to isolate the user space of the process.

## Software defined Infrastructure

Through the advent of large scale cloud endeavors and their offering of dynamic scaling services, the need for a way to automatically manage and partition the underlying infrastructure arose. This led to the development of software defined infrastructure, which is the process of abstracting the underlying infrastructure and managing it through software. The paradigm of software defined X is especially pronounced in the field of networking<sup>27</sup>. Software defined networking is used to manage the underlying network infrastructure, where the different underlying network devices are abstracted to a homogeneous network, which then gets managed through virtualization and software<sup>28</sup>. Large scale cloud providers like Amazon Web Services, Microsoft Azure and Google Cloud Platform all make use of software defined infrastructure to manage their underlying infrastructure.

---

<sup>24</sup>StackOverflow 2023

<sup>25</sup>N/A w.y.

<sup>26</sup>Docker 2023

<sup>27</sup>Xia et al. 2015, pp. 27–51

<sup>28</sup>Baur et al. 2015, pp. 95–101

## Large Scale Container Orchestration

These large scale software defined infrastructures are used in tandem and managed by large scale orchestrators like Openshift<sup>29</sup>, Openstack<sup>30</sup> and Kubernetes<sup>31</sup>. The orchestrators interface with the underlying infrastructure and manage the services running on top of it and provide a unified interface for the user to interact with. This makes it possible to largely run on any, even heterogeneous, infrastructure, while providing a highly flexible and scalable interface for the user to interact with. Make it especially useful for large scale cloud providers, as it allows them to provide a unified interface for their users, while still being able to use different underlying infrastructure.

## 3.2 High Performance Computing Frameworks

HPC is the process of using multiple computing nodes in collaboration to solve a problem in less time or with better accuracy than would be possible with a single node. The computational power provided by HPC is used in a wide variety of fields, ranging from weather forecasting to computational fluid dynamics and even in the field of machine learning. While HPC is as diverse as the problems it is used to solve, there are some common patterns that can be found in most HPC problems. Generally HPC problems can be divided into two categories, LCP and TCP problems, which are discussed in more detail in 3.2.2.

### 3.2.1 Loosely Coupled Problems

LCP also known in the industry as "embarrassingly parallel"<sup>32</sup> problems are problems that can be broken up into small independent tasks that can be executed in parallel. Problems like a monte carlo simulation or a matrix multiplication are good examples of LCP problems.

From the perspective of the user, LCP problems are the easiest to solve by using HPC, as they can be solved by simply running the same program multiple times with different input parameters across multiple nodes. Seen from the perspective of the HPC system, LCP problems trivial at best as they require very little communication between the different nodes.

The CC community has already seen the need in this aspect and has developed and integrated systems to solve these, such as containerized versions of MapReduce<sup>33</sup> or cloud native solutions like Pachyderm<sup>34</sup>

---

<sup>29</sup>Red Hat *OpenShift Enterprise Kubernetes Container Platform* 2023

<sup>30</sup>Openstack 2023

<sup>31</sup>*Production-Grade Container Orchestration* 2023

<sup>32</sup>Brown 2004

<sup>33</sup>Camacho-Rodríguez et al. 2019, pp. 1773–1786

<sup>34</sup>Pachyderm 2022a

### 3.2.2 Tightly Coupled Problems

In contrast to the highly paralelizable LCP problems, TCP problems are problems that can not be broken up into smaller independent tasks that can be executed in parallel, instead of working independently, each atomic task needs to communicate at least with one other task. A good example of a TCP problem are the n-body problems, where the position of each body is dependent on the position of all other bodies.

These kind of problems require a more sofisticated approach, as they need to share state across the multiple computing nodes. Sharing the necessary information between the nodes while maintaining performance at scale, has become its own sector in industry and science.

Communication protocols like Message Passing Interface (MPI)<sup>35</sup>, Open Shared Memory (OpenSHMEM)<sup>36</sup> have been developed to facilitate the communication between nodes and processes, but are still being continuously developed and optimized.

## 4 Creation of the Artifact

### 4.1 Initial Goals

As this project was first and foremost a project, designed to interactively explore the problem space from the perspective of the HPC community, all the while being contained by business requirements and time constraints, the initial goals of this project were very broad and open-ended. At first the initial goal was simply to create a Proof of Concept (PoC) of a realistic workflow engine using the "Arkouda" project, in order to present the Customer with an easily graspable example of its capabilities.

While we are approaching the problem from the perspective of the HPC community, the intended end user of this tool are the data scientists and Subject-Matter Experts (SMEs) that are working with the HPC systems, and therefore the tool needs to be designed and selected with the fact in mind that the end user will most likely not be knowledgeable in the field of HPC or the underlying infrastructure.

In the first iteration of the project a preselection of possible Workflow management tools was given from the business side, with the option to increase the scope if the presented tools were not sufficient.

Therefore, the goals of the first iteration of this project was twofold, first to determine which, if any, of the presented tools were suitable for the task at hand, and to determine what would make an adequate PoC for the customer.

---

<sup>35</sup>Snir 1998

<sup>36</sup>Chapman et al. 2010, pp. 1–3

The following iterations are split into the three main aspects of the project and will be discussed in their own subsections. While these steps were happening concurrently, they each addressed a different aspect of the project and therefore underwent their own iterative processes.

## 4.2 Overall Structure

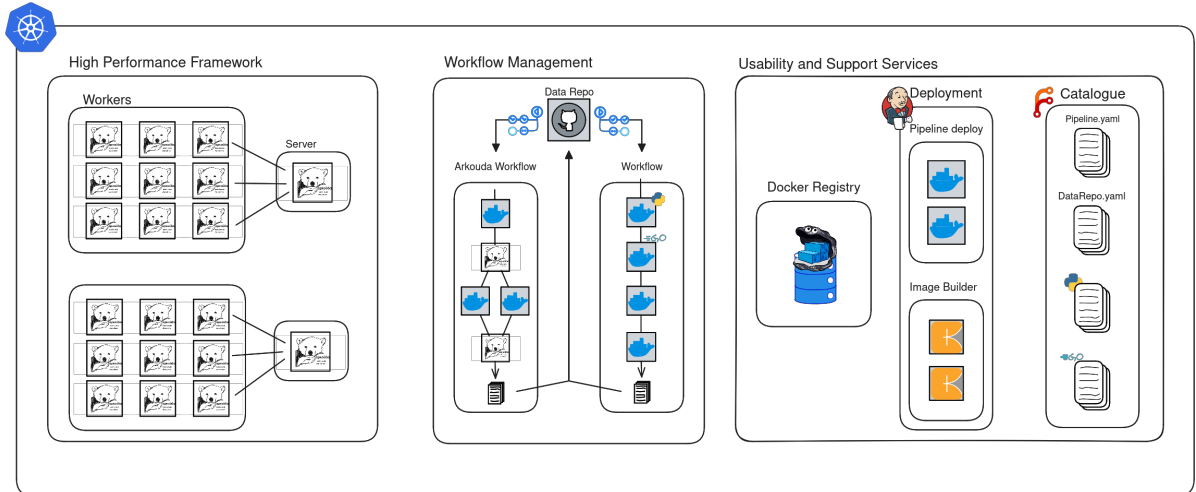


Abb. 6: Pachykouda high level infrastructure diagram

As can be seen in figure 6, the artifact is composed of three main components, the **Central Workflow Engine** which is responsible for the orchestration of the workflows (center) and interfaces directly with the underlying infrastructure, the **HPC Framework** which is responsible for the execution of TCP workloads (left) and the **Supplementary Services** which aim to improve the usability and accessibility for the end user (right).

All this is built on top of a hardware-agnostic k8s cluster, which is responsible for the orchestration of the different components and the underlying infrastructure.

## 4.3 Selection of Workflow Management Tools

As described in section 4.1, the first iteration of this project was to determine which, if any, of the presented tools were suitable for the task at hand. The following section will describe the process of selecting the tools and the criteria that were used to evaluate them. Because the time frame does not allow for a full integration and testing of all the presented tools in depth we will be using a decision-making framework to evaluate the tools, as described in the Methodologies 2.2, to determine which tools will be most suitable for an initial PoC and will serve as a good starting point for the project and future iterations.

- **Pachyderm:** A k8s based Workflow manager, written in go, which was recently acquired by Hewlett Packard Enterprise (HPE).
- **Argo:** A k8s based Workflow manager, written in go, which is a Cloud Native Computing Foundation (CNCF) project<sup>37</sup>.
- **Cloud Application Services Platform (CLASP):** An in-house developed workflow manager, written in Java, utilizing Serverlet to execute workflows<sup>38</sup>.
- **Snaplogic:** A commercial low-code/no-code workflow manager with a focus on data integration and data engineering<sup>39</sup>.

Given that it was possible to select projects outside the initial selection, the following projects also need to be considered:

- **Airflow:** A Python-based workflow manager under the CNCF umbrella, known for its easy-to-use interface and extensibility<sup>40</sup>.
- **Kubeflow:** A k8s-native platform for deploying, monitoring, and running ML workflows and experiments. Also CNCF project, streamlining ML operations alongside other Kubernetes resources<sup>41</sup>.
- **Knative:** An open-source k8s-based platform to build, deploy, and manage modern serverless workloads, simplifying the process of building cloud-native applications<sup>42</sup>.
- **Luigi:** An open-source Python module created by Spotify to build complex pipelines of batch jobs, handling dependency resolution, workflow management, and visualization seamlessly<sup>43</sup>.
- **Common Workflow Language (CWL):** An open-standard for describing analysis workflows and tools in a way that makes them portable and scalable across a variety of software and hardware environments, from workstations to cluster, cloud, and high-performance computing environments.

## Selection Criteria

Due to this extensive list of diverse tools, a set of criteria was established to determine which tool would be the most suitable for the task at hand. The following list of criteria was established to evaluate the tools:

---

<sup>37</sup>Argo 2023a

<sup>38</sup>Sayers, C. et al. 2015, pp. 1–60

<sup>39</sup>SnapLogic 2023

<sup>40</sup>Haines 2022, pp. 255–295

<sup>41</sup>Kubeflow 2023

<sup>42</sup>Knative 2023b

<sup>43</sup>Spotify 2023

- **Community, Support and Documentation:** Due to the limited time frame in which this project is being executed, the time to learn the intricacies of a given software is limited an important factor for this is therefore, that software needs to be adequately documented, and a support framework needs to be in place. Be it a community of users or a dedicated support team, the end users and the developers need to be able to rely on the software being maintained and updated as well as being able to find expert help in case of problems.
- **Maturity:** With the boom of AI and ML in recent years<sup>44</sup>, the number of tools and frameworks has exploded, and while this is a good thing it also means that a lot of these tools are still paving their way and continue to develop. While this is not necessarily a bad thing, it does mean that the tool may not be ready for production use and may not have the stability and reliability that is required for a production environment or are lacking in documentation and support.
- **Strategic alignment with HPE:** As this project is being developed within the context of HPE, it is important to consider the strategic alignment of the tool with HPE. HPE has is a large company with a diverse portfolio of products and services, and this project intersects with many parts of the company. Therefore, it is important to consider the strategic alignment of the tool with HPE and its products and services.
- **License:** While this PoC is not a commercial product in itself but rather an exploration of the problem space and a demonstration of what a final commercial product might be like, it is important to consider the licenses of the tools that are being used. Having to strip out a tool later on because of licensing issues would be a significant setback and therefore needs to be considered.
- **Cost:** Time is not the only constraint of this project, as the project is being developed within the context of HPE it is important to consider the cost of the tools that are being used.

### Weighting of the Criteria

An integral part of the SMART methodology is the weighting of the criteria, as described in section 2.2. In order to rank the criteria themselves, as they are quite hard to quantify, We will be using the weighting methodology as described in the SMART-ER methodology 2.2.3.

The first step of which is the ranking of the criteria from most important to least important.

1. **Community, Support & Docs** This also applies for the external support available to the development team, as if they are stuck, no development can proceed no matter the other factors.

---

<sup>44</sup>Forbes advisor 2023

2. **License** This criterion has to be weighted carefully, as a highly restrictive license might be a deal-breaker, but a license that is too permissive might conflict with the strategic alignment with HPE.
3. **Strategic alignment with HPE** As this is developed by and for HPE their requirements need to be considered as well.
4. **Cost** As this is a PoC and not a commercial product, the cost is not the highest priority as this will be of small scale and therefore the cost will be negligible in most cases.
5. **Maturity** While the maturity of the tool is important, as this is a PoC and not a commercial product, if the maturity of the tool does not impact the extensibility of the tool or the development process, it is not the highest priority.

As all these criteria are quite important, the weighting function selected for the criteria is the RS function, as described in section 2.2.3, as it does not rank the criteria too harshly. The lookup tables for the weighting function can be found in the appendix 16.

| Criteria                             | Weight |
|--------------------------------------|--------|
| Community, Support and Documentation | 0.3333 |
| License                              | 0.2667 |
| Strategic alignment with HPE         | 0.2000 |
| Maturity                             | 0.1333 |
| Cost                                 | 0.0667 |

Tab. 1: Weighting of the criteria

## Evaluation of the Tools

Now that we have established the criteria as well as their weighing, we can begin to evaluate the tools based on the criteria. Here we will be using a mix of methodologies, as some of these criteria can simply be indexed via analogous values, while others are of a more non-specific nature. The discussion of which values will be used on which weighing scale for the tools' comparison can be found in the appendix in the sections 1 - 5 1

| Criteria                  | Pachyderm | Argo | CLASP | Snaplogic |
|---------------------------|-----------|------|-------|-----------|
| Community, Support & Docs | 8.43      | 2.32 | 2.5   | 5.03      |
| Maturity                  | 7.86      | 5.2  | 2.5   | 5         |
| Strategic alignment       | 10        | 2.5  | 7.5   | 0         |
| License                   | 10        | 7.5  | 10    | 0         |
| Cost                      | 6         | 10   | 10    | 3         |

Tab. 2: Evaluation of the suggested tools

The following table shows the evaluation of the tools which were chosen for their relevance to the problem space, based on the criteria and the weighting of the criteria:

| Criteria                  | Airflow | Kubeflow | Knative | Luigi | CWL  |
|---------------------------|---------|----------|---------|-------|------|
| Community, Support & Docs | 10      | 2.25     | 0.74    | 2.29  | 0.22 |
| Maturity                  | 9.41    | 5.05     | 6.43    | 5.23  | 3.98 |
| Strategic alignment       | 2.5     | 2.5      | 2.5     | 2.5   | 2.5  |
| License                   | 7.5     | 7.5      | 7.5     | 7.5   | 7.5  |
| Cost                      | 10      | 10       | 10      | 10    | 10   |

Tab. 3: Evaluation of the additional tools

### Conclusion of the Selection Process

Now we can use these values and the weighting of the criteria to calculate the overall score of the tools, in order to find out which tool is the most suitable for the task at hand. The following table shows the overall score of the tools:

| Tool      | Score  |
|-----------|--------|
| Pachyderm | 8.9247 |
| Airflow   | 7.7546 |
| CLASP     | 6.0005 |
| Argo      | 4.6337 |
| Luigi     | 4.6277 |
| Kubeflow  | 4.5903 |
| Knative   | 4.2710 |
| CWL       | 3.7711 |
| Snaplogic | 2.5431 |

Tab. 4: Overall score of the tools

As can be seen in table 4, Pachyderm emerged as the most suitable tool, with the highest overall score of 8.9247. This suggests that Pachyderm aligns well with the project’s goals and the requirements of HPE, thanks to its strong community support, documentation, and strategic alignment with HPE’s objectives.

Airflow followed Pachyderm with a score of 7.7546, indicating its strength in community support and documentation, as well as its cost-effectiveness. Although Airflow did not score as high as Pachyderm, it remains a strong contender for use in future iterations of the project.

CLASP, an in-house developed tool, had a moderate score of 6.0005, which might make it suitable for initialPoC phases due to its alignment with the existing ecosystem and its no-cost advantage. However, its lower scores in other criteria might limit its potential for broader implementation.

Argo, Luigi, Kubeflow, and Knative scored in the middle range, suggesting they could be considered for specific use cases but may not be as well-rounded for the project’s current needs as Pachyderm or Airflow.



CWL's lower score of 3.7711 indicates it may not meet the project's requirements as effectively as the other tools evaluated. Its lower score in community support and strategic alignment with HPE's goals could be areas of concern.

Lastly, Snaplogic scored the lowest at 2.5431, suggesting that it may not be well-suited for this project's objectives, especially given its lack of a strong strategic alignment with HPE and lower scores in community support and documentation.

## 4.4 Implementation of the Artifact

This section will describe the iterative process of implementing the larger artifact and is broken up into 3 subsections. Since these steps were happening concurrently, they each address a different aspect of the project and therefore mostly underwent their own iterative processes.

### 4.4.1 Infrastructure

#### First iteration - Minikube

As the decision of the Workflow management tool was made, it was obvious that a dedicated k8s infrastructure was needed to run the tool<sup>45</sup>. The Pachyderm documentation gave two recommendations for setting up an initial development environment, preferably Docker Desktop or alternatively Minikube<sup>46</sup>. Due to the exclusive license of Docker-Desktop<sup>47</sup>, which prevents large companies free usage of the product<sup>48</sup> the choice fell on Minikube for an initial test setup.

In addition to the underlying k8s Pachyderm also needs an external S3 Storage Bucket for its Pachyderm File System (PFS) for which we used MinIO, a self-hostable S3 compliant object storage<sup>49</sup>, which was also based on recommendations by the Pachyderm documentation.

The persistent storage requirements for the Pachyderm itself was fulfilled by manually creating two Persistent Volume (PV)'s on the hosts local hard drive. Using the Helm packagemanager<sup>50</sup> for k8s the at that point, the newest version 2.6.4, was installed from the official Artifactory repository<sup>51</sup>.

The host system of this iteration was a single ProLiant DL385 Gen10 Plus running Ubuntu 22.04.3 LTS x86\_64. During the setup every step was diligently noted and put into a repository<sup>52</sup>, alongside the needed scripts. The instructions can be found in the appendix at 3.

---

<sup>45</sup>Pachyder 2023b

<sup>46</sup>Pachyder 2023a

<sup>47</sup>*Docker Terms of Service / Docker 2022*

<sup>48</sup>Docker 2021

<sup>49</sup>Inc 2023

<sup>50</sup>Helm 2023

<sup>51</sup>Pachyderm 2023b

<sup>52</sup>Eckert 2023a

### Learnings from the first iteration

The shortcomings of this naive first iteration became apparent very quickly, which was to be expected, as the goal of this iteration was to create a minimal working example to get a better understanding of the tooling and the underlying infrastructure.

The first and foremost issue were the limitations imposed by Minikubes' reliance on an Internal Virtual Machine (VM). During testing the inability to increase the resources of the VM on the fly became a significant bottleneck. At some point during the testing of 4.4.2 the VM was so overloaded that the installation was irreparably damaged which was seen as a sign to move on to the next iteration.

Another more subtle issue was the discrepancy between the experience a small scale k8s installation within Minikube and a large scale k8s cluster like the one that would be used in later steps of the project. Therefore, it was decided that a more realistic k8s cluster would be needed for the next iteration, which became the Heydar cluster.

### Second iteration - Heydar Cluster

Improving upon the shortcomings of the first iteration, the second iteration was based in the attempt to create a more realistic k8s cluster. To achieve this, 20 ProLiant DL360 Gen9 Servers, running Ubuntu 22.04.3 LTS x86\_64 were used to create a bare metal k8s cluster, using kubeadm as it provides deep integration with the underlying infrastructure<sup>53</sup>.

However a bare metal cluster also comes with its own set of challenges, as the cluster needs to be provisioned and configured manually. In order to automate this process, the Ansible automation tool was used to set up all the nodes in parallel and to ensure that all the nodes are in the same state. Ansible is a declarative tool which allows for the automation of the provisioning and configuration of the cluster<sup>54</sup>, by specifying the desired state of the cluster in a playbook and then applying it to the cluster. The Ansible playbook used for the setup of the cluster can be found in the projects repo<sup>55</sup>.

The application of this configuration on the cluster unknowingly caused conflict between the Ansible playbook and the maintenance scripts of the cluster as the Heydar machines. As k8s needs very specific configurations on the underlying infrastructure like the deactivation of swap space<sup>56</sup>.

This was resolved by consulting with the maintainer of the cluster and adjusting the Ansible playbook as well as the maintenance config for the cluster nodes accordingly, after we had identified the issue.

---

<sup>53</sup>Kubernetes 2023b

<sup>54</sup>Ansible 2023

<sup>55</sup>Eckerth 2023b

<sup>56</sup>Kubernetes 2023c

One important aspect of a production like cluster is the networking, as k8s does not natively manage communication on a cluster level, but instead relies on so called Container Network Interface (CNI)s to manage and abstract the underlying network infrastructure<sup>57</sup>.

Here we are spoiled for choice once again, as there are a multitude of different CNIs available, each with their own advantages and disadvantages. The Kubernetes documentation provides a non-exhaustive list of 17 different CNIs<sup>58</sup>, which all fulfill this essential task in different ways. As the needs regarding the network plugin were not very specific at this point, the choice fell on Calico, as surface level research showed that it was a popular choice for bare metal clusters<sup>59</sup>, provided security and enterprise support, as well having a wide range of features<sup>60</sup>. However Calico proved to be more difficult to set up than expected, after consulting with a colleague who set up a different cluster with Calico, it was decided to use Flannel as a CNI instead. Flannel turned out to be much easier to set up and configure, as it is a very lightweight CNI which is designed for bare metal clusters<sup>61</sup>, and foregoes the more advanced security features of Calico.

The Flannel configuration used for the cluster can be found in the project repo<sup>62</sup>, it is closely based on the example configuration provided by the Flannel documentation<sup>63</sup>.

### Learnings from the second iteration

The second iteration was a significant improvement over the first iteration, as it provided a much more realistic environment for the development of the artifact. This also came with its own set of challenges, as the bare metal cluster needed to be provisioned and configured manually, which was a significant time investment.

What became apparent very quickly was that the solution for the provisioning of the PV was nowhere near scalable, as it relies on the local hard drive of the host machine and therefore must host the container on the same machine as the PV which defeats the purpose of a multi node cluster in the first place. Therefore, a more scalable solution needs to be implemented for the next iteration. A possible solution could be the use of distributed storage solutions like Ceph<sup>64</sup> or GlusterFS<sup>65</sup> in combination with the Rook project<sup>66</sup>, which will need to be explored in future iterations.

As described in section 4.4.2 a service hosting Fabric Attached Memory (FAM) will be needed in future iterations as well.

---

<sup>57</sup>Kubernetes 2023a

<sup>58</sup>Kubernetes CNI Plugins 2023

<sup>59</sup>Bigelow 2023

<sup>60</sup>Mehndiratta 2023

<sup>61</sup>Flannel 2023a

<sup>62</sup>Eckerth 2023d

<sup>63</sup>Flannel 2023b

<sup>64</sup>CephIo 2023

<sup>65</sup>Gluster 2023

<sup>66</sup>Rook 2023

### 4.4.2 Tightly Coupled HPC Workloads

As described in section 3.2.2 TCP problems are a large part of the HPC world, but seem to lack native support in Pachyderm. Pachyderm, as it exists as of writing this thesis, is centralized around LCP problems, as it is designed to work with large amounts of data but with each so called "datum" being independent of each other. This is a very good fit for LCP problems, and ties into their concepts of data lineage, versioning and providence.

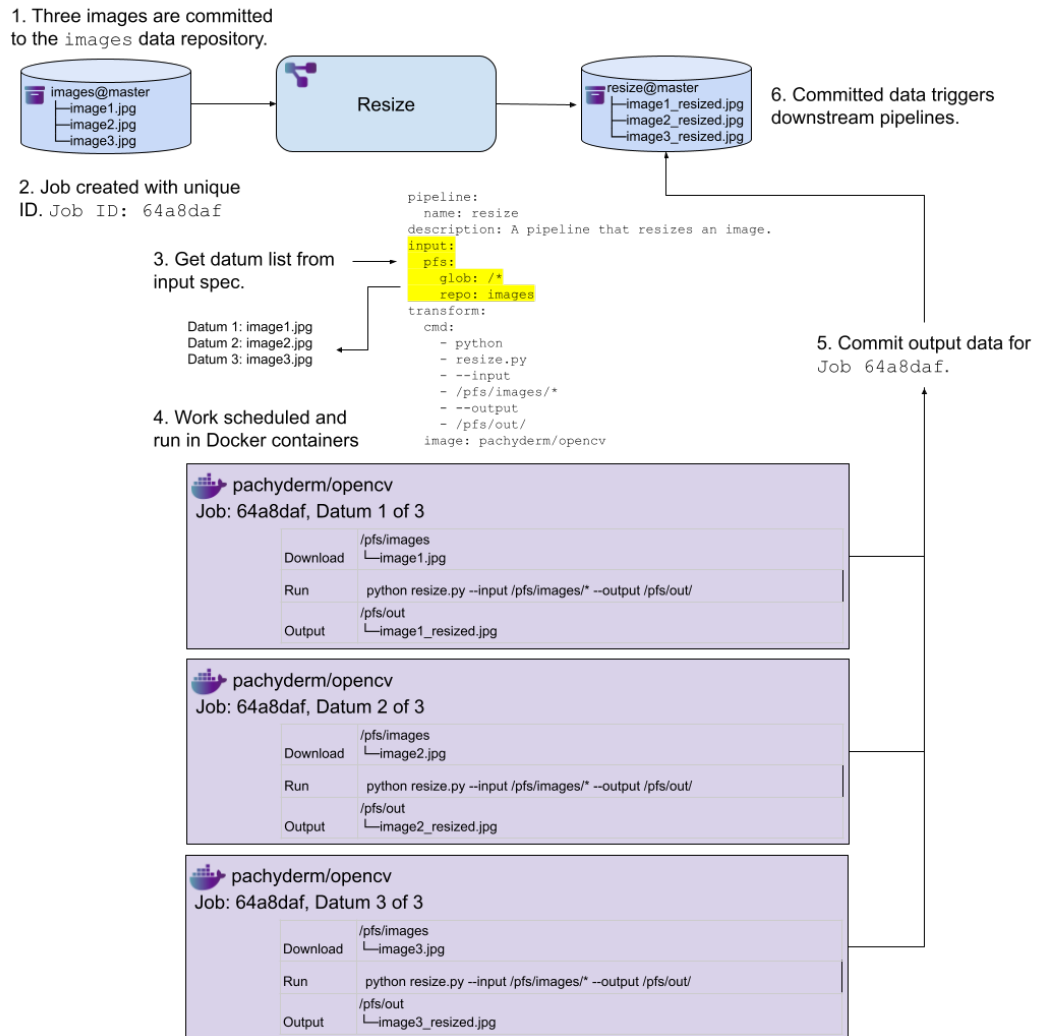


Abb. 7: Pachykouda datum distribution amongst workers <sup>67</sup>

Diagram 7 shows Pachyderm's approach to distribute their datums amongst workers, given an already defined pipeline. Once Data files are added to the input repository, Pachyderm will determine based on a glob pattern whether the files are relevant datums for the pipeline. If the newly added data fits the pattern each of the files will be supplied to its own instantiation of a worker, all originating from the same image, which will then process the data concurrently

<sup>67</sup>Taken from: *Intro to Pipelines* 2023

and independently of each other. After the worker has finished its task, the resulting datums are then collected in their own repository of data. A more detailed swim lane diagram of this process can be found in the appendix at 16

This approach is very well suited for LCP problems, as the datums are independent of each other and can be processed in parallel without any issues. Although it is not well suited for Large TCP problems, if the computation of the data can not be split into distinct independent datum files, or the computation is reliant on the intercommunication of the datums. If the datasets are small enough, this does not really present a problem as one can simply take all the data into a single worker node and process it there. However as a single worker node can only utilize the resources of a single physical compute node, this does not scale well with the size of the dataset and defeats the purpose of a distributed system in the first place.

So our goal for this section is to find a way to enable Pachyderm to pool the entire resources of the cluster, in order to solve a TCP problem.

### First iteration - PachyKouda

As a first attempt to address this issue, it was decided that the integration of a TCP framework into Pachyderm on the container level would be the best approach. So the first iteration is based on the idea of a Pachyderm conforming client container, which is able to interface with an external TCP framework, which can handle the reception of the data, the distribution of the data amongst the workers, and the collection of the results to reintegrate them into the PFS.

The first iteration of this idea was called PachyKouda, as it was based on the Arkouda TCP framework<sup>68</sup>, which itself is a python binding for the Chapel programming language<sup>69</sup>.

For that step an Arkouda worker was installed bare metal on the head node of the Heydar cluster, in order to verify the feasibility of the idea, with the goal of moving the worker into the cluster in the next iteration.

The client container was based on the official User Datagram Protocol (UDP)-based build by the Arkouda team<sup>70</sup>. The container was then modified to be able to communicate with the Arkouda worker on the head node of the cluster, it can now send data to the worker and receive the results.

### Learnings from the first iteration

The first iteration was a total success, as it proved the feasibility of being able to use a client container to forward the data processing to an external Arkouda worker. As described earlier,

---

<sup>68</sup>Bears-R-Us 2023a

<sup>69</sup>Hewlett Packard Enterprise Development LP 2023

<sup>70</sup>Bears-R-Us 2023b

the goal of the next iteration is to move the Arkouda worker into the cluster, in order to be able to utilize the full resources of the cluster.

## Second iteration - Kymera

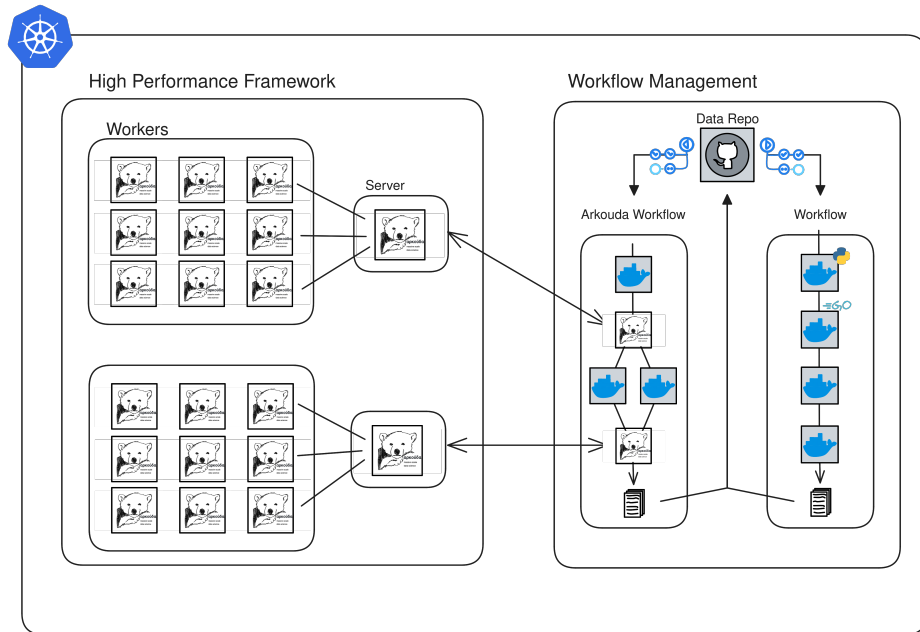


Abb. 8: Arkouda workers on the Heydar cluster

Diagram 8 above shows a high level overview of how the workers interface with the client container in the workflow. The Arkouda container, which is part of the workflow is still the same as in the first iteration, but now instead of interfacing with an external worker it is interfacing with a worker swarm hosted across the cluster.

The Swarm is split into two parts: one central Arkouda server, facilitating the communication between the client container and the workers, and the workers, also called locales themselves. The locales and the server are based on the helm charts provided by the Arkouda-Contrip repo<sup>71</sup>,

A detailed walk through the setup of the Role Based Access Control (RBAC), Secrets and deployments for the Heydar Cluster can be found in the project repo<sup>72</sup> which in turn is based on the official Arkouda documentation<sup>73</sup>.

---

<sup>71</sup>Bears-R-Us 2023c

<sup>72</sup>Eckerth 2023f

<sup>73</sup>Bears-R-Us 2023b

### Learnings from the second iteration

As Arkouda does not currently provide multi-tenancy of their server, meaning that they can only be connected to a single client at a time, if multiple pipelines need to solve a TCP at the same time, they would not be able to share the same worker swarm. Instead, they would need to spawn their own worker swarm.

Another issue is that there are currently going through the standard pod to pod communication configuration of flanne. This means that the entire traffic between the client container and the Arkouda server as well as the traffic between the worker,s is all happening over emulated overlay network.This enables the containers on the different nodes to communicate with each other as if they where on the same network, no matter of the actual infrastructure below it. The communication protocol of the Arkouda servers is UDP based Global Address Space Networking (GASNet), which provides the Remote Direct Memory Access (RDMA) needed for the Arkouda framework to work, but this incurs a significant overhead in the form of the encapsulation of the UDP packets into TCP packets.

Also, the containers are currently not compatible with the OpenFAM project<sup>74</sup>, which is being developed as an integration to Arkouda and Chapel by the Hewlett Packard Systems Architecture Lab<sup>75</sup>, it extends the Arkouda framework with the ability to use FAM as banks of RDMA enabled memory, which can be accessed by the Arkouda workers. This would prove to be a significant improvement as it has the potential to reduce the overall overhead of the communication<sup>76</sup> amongst the workers as well as to the server, by cutting down the overall amount of network traffic.

The pachyderm platform itself might also benefit from the integration of FAM, as it could be used to store the datums in the PFS, providing the running pipeline processes with a much faster access to the data.

### Third iteration - FAM

While significant efforts have already been made to successfully integrate Arkouda and FAM, these have so far been focussing on bare metal installations, for that reason, in order to integrate the FAM enabled Arkouda working from within a containerized environment, the tools would need to be custom recompiled matching the new environment. Therefore, we needed to:

1. Compile OpenFAM in the Container.
2. Compile custom Chapel in the Container with OpenFAM.
3. Compile custom Arkouda in the Container with the OpenFAM enabled Chapel.

---

<sup>74</sup>Keeton/Singhal/Raymond 2019, pp. 70–89

<sup>75</sup>Byrne et al. 2023

<sup>76</sup>Chou et al. 2019, pp. 16–24

4. Rebuild the Arkouda container with the new Arkouda binary.
5. Rewrite the k8s deployment to make use of OpenFAM.

This section was quite challenging as it required a deep understanding of the PoC implementations of the OpenFAM, Arkouda and Chapel projects and was cut short by the time constraints of the project and was therefore not brought to a successful conclusion. The current state of the project can be found in the PoC repository<sup>77</sup>.

But this showed us that there is a lot of potential in the integration of FAM into the container based HPC world, as it could provide a significant performance boost to the overall system and should be explored further in future iterations.

#### 4.4.3 Supplementary Services

As the other branches of the prototyping were happening, the need for additional services and infrastructure arose to support the development of the prototype as well as to increase the general usability of the prototype. This section will especially describe the services which helped to make this prototype a more complete solution.

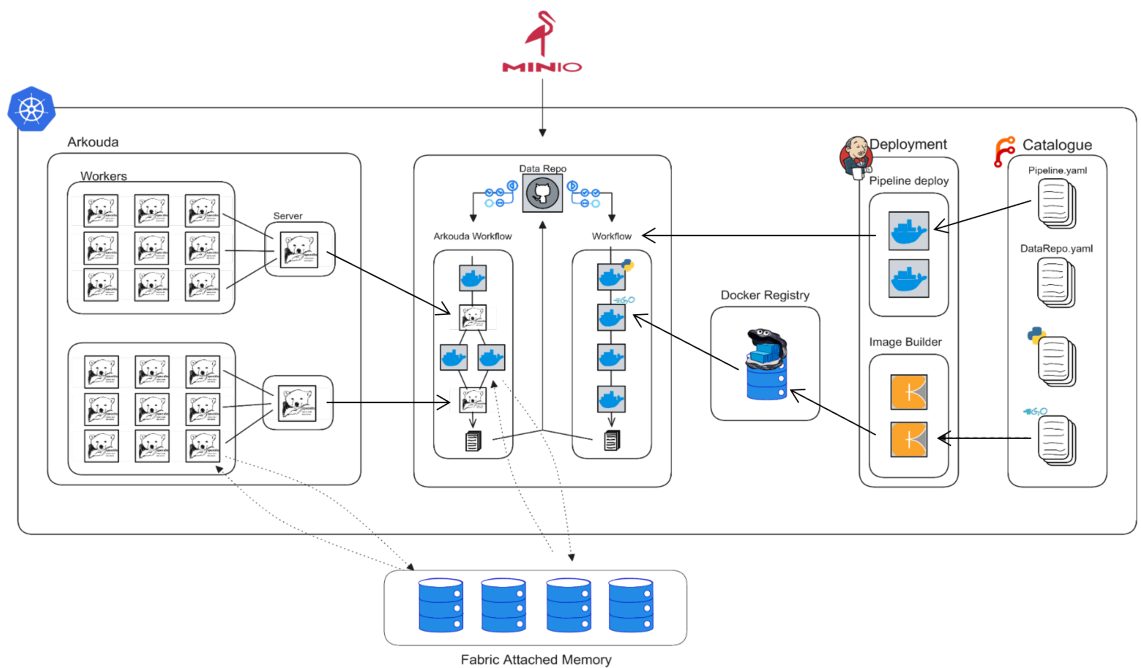


Abb. 9: Pachyderm High-Level Architecture

<sup>77</sup>Eckerth 2023g



### Docker Registry

One thing that was quite apparent from the get go was the need for a central docker registry. As Pachyderm does not manage the docker images itself, but relies on the user to provide them somehow externally.

During the first iterations when the development was being done on Minikube as described in 4.4.1, the internal registry of the node was enough. But as soon as we moved over to the Heydar system keeping the Hosts internal registries in sync was of course not feasible, Therefore we added a private docker registry to the cluster<sup>78</sup>. The deployment config is based on the official docker registry helm chart<sup>79</sup> and can be found in the projects repo<sup>80</sup>.

### Forgejo Catalogue

As this should be the PoC of an end user tool, we should also look into usability features and based on previous experience of the team, the need for a catalogue of previously developed pipelines and processing code was identified. The idea was to create something similar to the CLASP catalogue<sup>81</sup> but for the Pachyderm ecosystem. This means that users can share, search and deploy workflows from a central catalog, without having to worry about the underlying infrastructure.

Having HPC software in a completely contained and versioned system directly addresses many of the original problem statements, described in 1.2, especially the problem of reproducibility, environment management, and the lack of portability.

To start off we, installed a simple github-like interface for the catalogue, called Forgejo<sup>82</sup>, which is a fork of the Free and Open Source Software (FOSS) project Gitea<sup>83</sup>, being maintained by the Non-Governmental Organization (NPO) Codeberg e.V.<sup>84</sup>. The installation was done using their official helm chart<sup>85</sup> and can be found at in the project repo<sup>86</sup>.

### Jenkins

Now that we have a place to hold and version the code files and a place to hold the resulting docker images, we need a way to build and deploy them automatically. For that purpose an installation of Jenkins was added to the Cluster. Jenkins is a FOSS CI/CD tool, which is widely

---

<sup>78</sup>Kumar 2020

<sup>79</sup>Phntom 2023

<sup>80</sup>Eckerth 2023c

<sup>81</sup>Sayers, Craig et al. 2015

<sup>82</sup>Forgejo 2023a

<sup>83</sup>Gitea 2023

<sup>84</sup>Codeberg 2023

<sup>85</sup>Forgejo 2023b

<sup>86</sup>Eckerth 2023e

used in industry<sup>87</sup>. It enables us to execute arbitrary code in an controlled environment based on triggers or schedules. We will be using it to automatically build and deploy docker images whenever code is pushed to the Forgejo catalogue, as well as automatically deploying pipeline scripts to Pachyderm.

By using this we ensure that the code and the docker images are always in sync and that the user does not have to worry about building and deploying the images manually or interacting with the underlying Cluster. We also extend the factor of provenance which was limited to the data itself, to the containers code and pipeline spec as well, now having total oversight into which input begets what output. This was achieved in multiple steps as this turned out to be quite an involved process.

First off was of course the general installation of the project into the cluster based on the image provided by Jenkins<sup>88</sup>. Then we have to integrate Jenkins into the the Kubernetes cluster as well, as we want it to be able to spawn its own worker pods within the cluster to handle the building of the containers and pushing of the pipelines. Also an integration with the Forgejo catalogue was needed, so that Jenkins would be informed about new commits to the catalogue and could start the building process.

The installation instructions, RBAC and the configuration of the Jenkins installation can be found at 4.

### CI/CD Pipeline

Now that we have a Jenkins installation which is able to spawn its own worker pods on the cluster and have set up system wide Webhooks for the Forgejo catalogue, which will inform Jenkins about every new commit to any repository on the catalogue, we will need to develop a pipeline which will tell Jenkins what to do if a new commit is detected.

As the Pachyderm team say themselves, "*[/.../ users with limited experience with containerization, cloud computing, and distributed systems may find it challenging to use Pachyderm effectively.*"<sup>89</sup> Unfortunate many SMEs are more focused on their domain specific knowledge and usually only have a limited understanding of the underlying infrastructure, an effect which has already been noticed in classical HPC<sup>90</sup>, which is likely more pronounced in a field which has only recently started to gain traction in the wider scientific community.

Therefore providing an way to deploy code and pipelines while minimizing the interaction with the underlying infrastructure is a key feature of this part of the prototype. The goal is to create a pipeline which can take a regular software project form the Forgejo catalogue, make reasonable

---

<sup>87</sup>6sense 2023

<sup>88</sup>Jenkins 2023a

<sup>89</sup>Pachyderm 2023a

<sup>90</sup>Shenoi/Shah/Joshi 2019, pp. 64–70

assumptions about the project structure and build the required docker images and deploy the pipeline to Pachyderm.

A completely functional Pachyderm project, consisting of a  $N$  processing steps requires the following components:

- $N$  code files, one for each processing step
- $N$  Dockerfiles, one for each processing step
- $N$  pipeline specifications, one for each processing step
- At least one description file for the data repository(s)
- If needed supporting files, which should be included in one ore more image, but should not get their own.

However as we want to minimize the interaction with the underlying infrastructure, we want to minimize the amount of input the user has to provide, without restricting the knowing users. In order to achieve this, the pipeline first detects the existence of the above components. We then try to make a reasonable assumption about the structure of the project and the relation between the components, considering: directory structure, naming conventions, and static code analysis. The pipeline will then try to fill in gaps in the project structure, like missing Dockerfile by using default values and insights gained form the previous steps. The pipeline itself is written in a mix of Groovy, Python and bash, a flowchart describing the process in more detail can be found in the appendix at 17 its code can be found in the projects repo<sup>91</sup>.

There was also a custom container image created, extending official Jenkins worker image<sup>92</sup>, enabling us to build docker images from within a running container in the cluster by using a dedicated kaniko sidecar container<sup>93</sup> which enables us to build docker images without having to run Docker in Docker, which is not recommended<sup>94</sup>. The image also contains the necessary tools to interact with the Pahcyderm API and to run the Python scripts which are used to analyze the project structure and do t e code generation, it can be found in the project Repo<sup>95</sup>.

While this pipeline does technically work, it is still quite finicky and is missing many edgcases and does need further development to be able to handle more thant the most basic projects. Even though its development was cut short it does show the potential of this approach and how it could be used to make the Pachyderm ecosystem more accessible to users which are not as familiar with the underlying infrastructure, while still allowing more experienced users to interact with the system directly.

---

<sup>91</sup>**Pipeline**

<sup>92</sup>Jenkins 2023b

<sup>93</sup>Google 2023

<sup>94</sup>Petazzoni 2023

<sup>95</sup>**PipelineImage**

## 4.5 Evaluation of the Artifact

The original problems described in 1.2 were seven-fold, and were addressed in the following way:

- **Workload Resilience and Fault Tolerance in HPC:** A problem which is typically addressed by the HPC community by using a combination of checkpointing and job scheduling<sup>96</sup> is now being directly addressed via the inclusion of Pachyderm. Since Pachyderm isolates each of the processes into their own container, and tracks each of the steps individually it can easily restart a failed step, or even a failed job, without having to restart the entire workflow. While this only works for the standard LCP workloads, its provenance features reduce the data loss, should a TCP workload fail.
- **Environment/Package Management in HPC:** Replacing the classical HPC package management solutions with a containerized approach, simplifies the deployment of code from the users' perspective massively, as they have almost complete control of the environment their code is going to run in, all the while giving the administrators the ease of mind that the code is not going to interfere with the rest of the system.
- **Probability issues with HPC:** Same goes for the probability issues, as the containerized approach allows for a much more fine-grained control of the environment, users can rapidly iterate and test their code on their local machines, before deploying it to the HPC system.
- **Scalability issues with HPC:** Scalability is one of the main features of k8s, and therefore of Pachyderm, a cluster can easily be scaled up or down, depending on the current workload. Many cloud providers even offer hosted k8s clusters, which can be scaled up or down on demand, and therefore allow for a very flexible approach to the problem. While in classical HPC systems, the cluster is usually fixed in size, and therefore the user has to wait for the next available slot.
- **Interconnected Problem-Solving in CC** This one was one of the problems which was not directly addressed by Pachyderm or Kubernetes directly, in order to solve this problem, Arkouda was containerized and made usable for Pachyderm workloads. As of right now, the layers of network abstractions and the lack of OpenFAM support have a negative impact on the performance of the TCP workloads, but successfully proves the concept of interconnected problem-solving on a CC system.
- **Provenance and Versioning:** Combining the advantages the Pachyderm File System with the completely CI/CD based approach to the deployment of the workflows, allows a tracking of each and every part that goes into each and every step of the workflows.

While this project does not present a complete solution to all the problems, it does present a viable path forward for a more modern approach to HPC. The combination of HPC and

---

<sup>96</sup>Jin et al. 2010, pp. 525–534

Pachyderm allows for a much more flexible approach to the problem and with future work and low level driver support and usability features like the Jenkins Pipeline, as well as a well maintained ecosystem of pipeline steps which can be used to build more complex workflows and reutilize existing code developed by fellow researchers, this approach could be a significant improvement over the current state of the art.

Unfortunately the time was cut short before the project could be fully completed and therefore some goals like the integration of OpenFAM, the switch to a low abstraction CNI and especially multi parameter performance testing could not be completed in time, it is apparent that the integration of these would bring the project much closer to the performance of the classical HPC systems, while still maintaining the flexibility and ease of use of the containerized approach, and therefore should be considered for future work.

## 5 Summary and Outlook

### 5.1 Summary

By way of this project paper we have assessed and described the state of the art for current technologies in the areas of containerized software, container orchestration and how those tie in with Software defined infrastructure. We have done the same for the state of the art in the solution of complex problems found in the field of and solved with HPC, namely the two larger classes of problems, LCP and TCP problems, and what strategies are usually employed to solve those problems.

We have then addressed the problems identified by the problem statement with the statement of an initial goal, and structured an intervention to solve those problems by way of the creation of a prototype. Before the iterative process of prototyping could begin we had to make a non-trivial decision on the choice of a container orchestrator/ workflow manager. For which we defined, discussed and weighted the selection criteria, and then evaluated the available options against those criteria, and We found our best fit through the employment of the SMART-ER method, which was the k8s-based orchestrator named Pachyderm.

The iterative process of prototyping was split into three main areas of focus, the infrastructure, the solution to TCPs and the integration of a complete CI/CD pipeline. Each of these areas was given at least two iterations and the results of each iteration were documented, evaluated and recommendations for future iterations were made.

The final results of the prototype were then evaluated against the initial goal and the problems identified in the problem statement, and the resulting artifact was found to be a valid form of intervention to solve the problems identified in the problem statement, while still being limited in scope and applicability, as was to be expected from a non production prototype.

## 5.2 Outlook

As already discussed, the limiting factor of this project was the available time, which left many avenues to still be explored both in respect to the prototype itself, and how it can be further extended to explore avenues not yet addressed by this project.

Of most importance in this case would be a further exploration of an efficient solution to TCP problems, as the current solution does technically work, but is lacking the efficiencies to be a viable solution for large scale problems. An initial starting point for this could be the network stack in between the containers and the underlying network infrastructure.

# Appendix

## Appendix Index

|                                                                 |    |
|-----------------------------------------------------------------|----|
| Appendix 1 Discussion of Tool Evaluation and Weighing . . . . . | 32 |
| Appendix 1/1 Community, Support & Docs . . . . .                | 32 |
| Appendix 1/2 License . . . . .                                  | 33 |
| Appendix 1/3 Strategic alignment . . . . .                      | 35 |
| Appendix 1/4 Maturity . . . . .                                 | 36 |
| Appendix 1/5 Cost . . . . .                                     | 37 |
| Appendix 1/6 Lookup table weighing functions . . . . .          | 38 |
| Appendix 1/7 Pipeline Communication Swim Lane Diagram . . . . . | 40 |
| Appendix 2 Diagrams . . . . .                                   | 41 |
| Appendix 2/1 Jenkins Pipeline Diagram . . . . .                 | 41 |
| Appendix 3 Minikube installation instructions . . . . .         | 43 |
| Appendix 4 CI/CD installation instructions . . . . .            | 50 |

## Appendix 1: Discussion of Tool Evaluation and Weighing

### Appendix 1/1: Community, Support & Docs

This section assesses the level of external support provided for each project. To evaluate this support, we will focus on three distinct aspects and combine them into a single score. Firstly, we will examine the size of the community, as a substantial community often indicates project maturity and the availability of extensive support. As proxies for community size, we will consider two central metrics: the number of stars on GitHub and the quantity of questions on Stack Overflow.

Tab. 5: Community Support Data<sup>97</sup>

| Project   | GitHub Stars | Stack Overflow Questions |
|-----------|--------------|--------------------------|
| Pachyderm | 6,000        | 6                        |
| Argo      | 14,500       | 136                      |
| Clasp     | 0            | 0                        |
| Snaplogic | 0            | 57                       |
| Airflow   | 32,200       | 10,218                   |
| Kubeflow  | 13,100       | 434                      |
| Knative   | 4,100        | 204                      |
| Luigi     | 16,900       | 346                      |
| CWL       | 1,400        | 6                        |

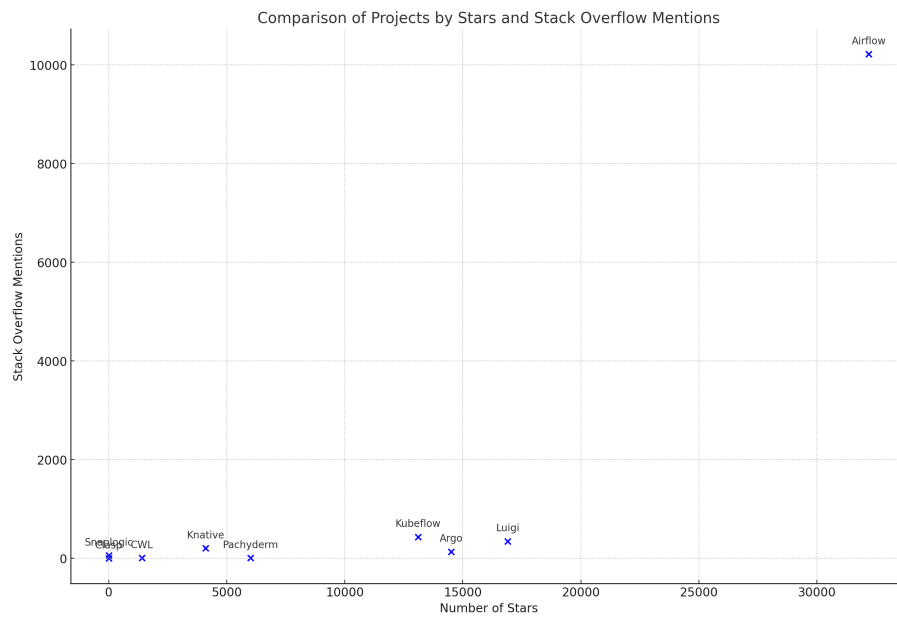


Abb. 10: Stars and Stack Overflow Questions Comparison

To gauge the level of support and community engagement surrounding these projects, we have devised a composite score that normalizes and combines the GitHub stars and Stack Overflow questions metrics. The calculation of this score involves the following methodology:

<sup>97</sup>Data taken from GitHub and Stack Overflow on 08.11.2023



Each project is represented as a point  $P_i = (x_i, y_i)$  in a two-dimensional space, with  $x_i$  and  $y_i$  being the number of GitHub stars and Stack Overflow questions, respectively, for the  $i$ -th project. The composite score  $S_i$  for each project is computed by normalizing these values to a scale of 0-10 and then taking their average.

Additionally, we acknowledge that some commercial tools, as well as certain open-source projects, offer enterprise support, reducing the reliance on the community for assistance. Similarly, projects developed in-house often have access to the original development team for support. Therefore, we will apply a flat bonus of 5 points to the scores of projects offering enterprise support and a flat bonus of 2.5 points to projects developed in-house.

$$S_i = \frac{1}{2} \left( \frac{x_i - \min(x)}{\max(x) - \min(x)} \times 10 + \frac{y_i - \min(y)}{\max(y) - \min(y)} \times 10 \right) + B_i$$

Here,  $\min(x)$ ,  $\max(x)$ ,  $\min(y)$ , and  $\max(y)$  represent the minimum and maximum values of GitHub stars and Stack Overflow questions across all projects, respectively. The final scores  $S_i$ , along with the respective bonuses  $B_i$ , provide a comprehensive metric for comparing project popularity, community engagement, and the availability of additional support options, all on the same scale.

| Project   | Composite Score | Enterprise Bonus | Inhouse Bonus | Final Score |
|-----------|-----------------|------------------|---------------|-------------|
| Airflow   | 10.00           | 0                | 0             | 10.00       |
| Pachyderm | 0.93            | 5                | 2.5           | 8.43        |
| Snaplogic | 0.03            | 5                | 0             | 5.03        |
| Luigi     | 2.79            | 0                | 0             | 2.79        |
| Clasp     | 0.00            | 0                | 2.5           | 2.5         |
| Argo      | 2.32            | 0                | 0             | 2.32        |
| Kubeflow  | 2.25            | 0                | 0             | 2.25        |
| Knative   | 0.74            | 0                | 0             | 0.74        |
| CWL       | 0.22            | 0                | 0             | 0.22        |

Tab. 6: Composite scores of Workflow managers, sorted by final score

## Appendix 1/2: License

As discussed in section 4.3 the tools in consideration should not be too restrictive. To evaluate the criteria we will employ a 4 bucket system:

- **Ideal Situation (Score: 10):** This refers to cases where either the tool is in the public domain (and therefore not subject to copyright restrictions) or where our organization possesses a direct ownership or significant influence over the licensing terms. This situation provides the most flexibility, allowing for extensive modification, redistribution, and proprietary use without concern for licensing infringements.

- **Permissive License (Score: 7.5):** Tools under licenses like MIT, BSD, or Apache 2.0 fall into this category. These licenses are highly permissive and generally allow for broad freedom, including modification, distribution, and private use, with minimal restrictions, often limited to liability and warranty.
- **Restrictive or Reciprocal Licenses (Score: 2.5):** Licenses such as the GPL or AGPL are more restrictive, requiring any changes to be open-sourced or contributions to be made back to the community. These “copyleft” licenses can be problematic in proprietary settings where modifications or integrations need to remain confidential.
- **Unacceptable Licenses (Score: 0):** This includes licenses that impose burdensome conditions or high costs, proprietary software where the source code is unavailable, or situations where the licensing terms make it impractical to use within our projects. For instance, licenses that mandate the purchase of additional software, restrict certain types of use, or pose potential legal risks would fall into this category.

Now we will evaluate the licenses of the tools in question, and assign them a score based on the above criteria.

- **Pachyderm** The licensing model of Pachyderm follows a model which has similarities with the "Open Core model"<sup>98</sup>. Which means that while the core functionalities are published as the "COMMUNITY EDITION" with a permissive source-available License (Apache License 2.0)<sup>99</sup>. Functionality like Single Sign On (SSO) or the ability to create more than 16 pipelines are part of a different distribution under a Commercial License.

But in our case this is of no concern, as the startup behind the Pachyderm software, including its Intellectual Property (IP) was acquired by HPE. Giving us a free hand to modify without needing to worry.

- **Argo** Argo’s adoption of the Apache License 2.0<sup>100</sup> aligns with common practices for open-source projects, affording users considerable freedom. This permissive license simplifies the use, modification, and redistribution of the software, an aspect that’s particularly beneficial for collaborative development or integration into proprietary software. Given our requirements and operational context, this offers us the flexibility needed for adaptation and potential enhancements without stringent restrictions, streamlining any developmental efforts we undertake with Argo.
- **CLASP** is not a published software and therefore not under any specific license. But similar considerations as the ones of Pachyderm apply here as well, as it is an internal project the IP also completely belongs to HPE

---

<sup>98</sup>Pachyderm 2022b

<sup>99</sup>Pachyderm 2023c

<sup>100</sup>Argo 2023b

- **Snaplogic** is an entirely commercial product which does not provide insight into nor the right to modify their Software<sup>101</sup>. But as they might agree this is not a total knockout criterion for this entire project, but in regard to the licensing it will be weighted with 0.
- **Airflow** is licensed under the Apache License 2.0.<sup>102</sup>
- **Kubeflow** is licensed under the Apache License 2.0.<sup>103</sup>
- **Knative** is licensed under the Apache License 2.0.<sup>104</sup>
- **Luigi** is licensed under the Apache License 2.0.<sup>105</sup>
- **CWL** is licensed under the Apache License 2.0.<sup>106</sup>

### Appendix 1/3: Strategic alignment

When considering the strategic alignment within the context of HPE the goal is to ensure that the tool can be used in combination with or as extension to the Products HPE is offering. Tools that have been developed by the company in the first place or have been acquired by HPE must be of a high strategic value to the company, as otherwise they would not have been developed or acquired in the first place.

This is especially true for Pachyderm, as it has been acquired in 2023<sup>107</sup> as part of their realignment towards hybrid cloud provider. The eight years since the development of CLASP in 2015<sup>108</sup> the priorities of somewhat shifted, its concept is still relevant and was originally developed by the labs to fill a strategic gap.

While some of these tools might have been used in other projects, been contributed to by HPE employees or are supported on the Ezmeral platform, they are not part of the core product portfolio and therefore not of strategic importance.

One thing that would have to be considered is the way one would make oneself dependant of the company, as HPE would have to enter into negotiations with the company behind Snaplogic, this state of dependency and the need to find a compromise would be a strategic disadvantage.

Therefore the following scores have been assigned:

---

<sup>101</sup> *SnapLogic – Master Subscription Agreement 2023*

<sup>102</sup> *Apache 2023*

<sup>103</sup> *Kubeflow 2023*

<sup>104</sup> *Knative 2023a*

<sup>105</sup> *Luigi/LICENSE at Master · Spotify/Luigi 2023*

<sup>106</sup> *Common-Workflow-Language 2023*

<sup>107</sup> *Enterprise 2023*

<sup>108</sup> *Sayers, Craig et al. 2015*

Tab. 7: Strategic Alignment Scores

| Project   | Strategic Alignment |
|-----------|---------------------|
| Pachyderm | 10                  |
| CLASP     | 7.5                 |
| Argocd    | 2.5                 |
| Snaplogic | 0                   |
| Airflow   | 2.5                 |
| Kubeflow  | 2.5                 |
| Knative   | 2.5                 |
| Luigi     | 2.5                 |
| CWL       | 2.5                 |

## Appendix 1/4: Maturity

The evaluation of the tools is centrally based on their open source repositories, as this is the most accessible and transparent source of information. In order to evaluate the maturity of the tools, we will look at the age, the number of commits, the ratio of commits to age, the number of closed issues and the ratio of closed issues to open issues. We will then combine these metrics into a single score, which will be used to evaluate the maturity of the tools. Products like Snaplogic, which are not open source, as they are a commercial product it is to be expected that they are somewhat mature, as they are already being used in production environments and are used by many companies.

| Project   | Age (Days since 08.11.2023) | Commits | Open Issues | Closed Issues |
|-----------|-----------------------------|---------|-------------|---------------|
| Pachyderm | 3352                        | 22143   | 889         | 2387          |
| Argo      | 2098                        | 6047    | 2693        | 4161          |
| Airflow   | 3131                        | 22026   | 762         | 7467          |
| Kubeflow  | 2168                        | 2514    | 178         | 3615          |
| Knative   | 2113                        | 8531    | 202         | 4332          |
| Luigi     | 4066                        | 4088    | 93          | 879           |
| CWL       | 2954                        | 4512    | 421         | 358           |

Tab. 8: Maturity Data<sup>109</sup>

How these values compare to each other in normalized form is shown in the following graphs:

<sup>109</sup>Data taken from GitHub on 08.11.2023

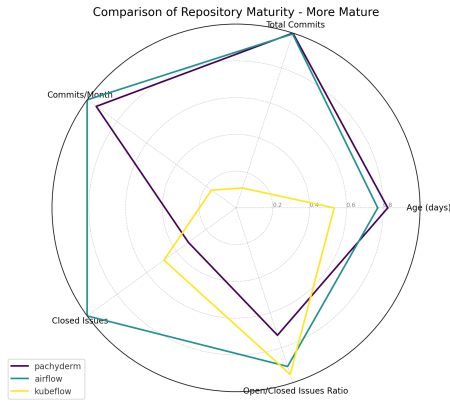


Abb. 11: Maturity Metric comparison

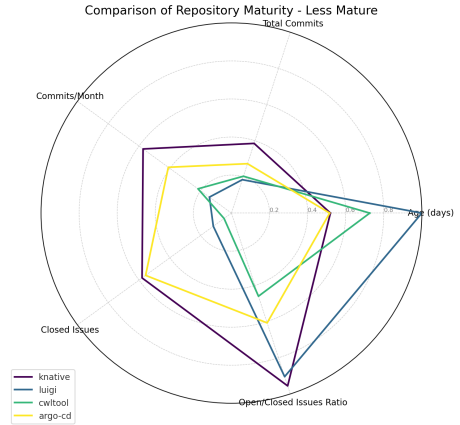


Abb. 12: Maturity Metric comparison

The combined maturity score for each tool is calculated as follows:

- Age (days): Current Date - Created Date
- Commits/Month: Total Commits / (Age in months)
- Open/Closed Issues Ratio: Closed Issues / (Open Issues + Closed Issues)
- Normalized Metrics: Each metric is normalized by dividing by the maximum value among all repositories.
- Combined Maturity Score: Average of normalized metrics, scaled to a 1-10 range as  $1 + (9 \times \text{average})$ .

Tab. 9: Maturity Scores

| Project   | Maturity Score |
|-----------|----------------|
| Pachyderm | 7.86           |
| Argo      | 5.2            |
| Airflow   | 9.41           |
| Kubeflow  | 5.05           |
| Knative   | 6.43           |
| Luigi     | 5.23           |
| CWL       | 3.98           |

## Appendix 1/5: Cost

This section aims to compare the relative cost of the products in relation to each other. We previously factored in the enterprise features, so when enterprise support is available and applicable we will take this into consideration. Here we have three categories of products first those which are completely free and without any enterprise support, secondly those which are free but offer

enterprise support and lastly those which operate on a subscription basis. To address various cost models, we classify the tools into the following categories:

1. **Open-source and Community-Supported Tools (No Cost):** Free tools, typically open-source, relying on a community for maintenance.
2. **Open-source with Optional Enterprise Support (Variable Cost):** Open-source with paid enterprise support offering advanced features by the developing team.
3. **Commercial Products (Subscription-Based Cost):** Tools that require ongoing payment for use depending on the level of service and feature access.

Tab. 10: Scores Based on Cost

| Project   | Cost Score |
|-----------|------------|
| Airflow   | 10         |
| Pachyderm | 6          |
| Snaplogic | 3          |
| Luigi     | 10         |
| CLASP     | 10         |
| Argo      | 10         |
| Kubeflow  | 10         |
| Knative   | 10         |
| CWL       | 10         |

## Appendix 1/6: Lookup table weighing functions

| Rank | Attributes |        |        |        |        |        |        |        |        |
|------|------------|--------|--------|--------|--------|--------|--------|--------|--------|
|      | 2          | 3      | 4      | 5      | 6      | 7      | 8      | 9      | 10     |
| 1    | 0.7500     | 0.6111 | 0.5208 | 0.4567 | 0.4083 | 0.3704 | 0.3397 | 0.3143 | 0.2929 |
| 2    | 0.2500     | 0.2778 | 0.2708 | 0.2567 | 0.2417 | 0.2276 | 0.2147 | 0.2032 | 0.1929 |
| 3    |            | 0.1111 | 0.1458 | 0.1567 | 0.1583 | 0.1561 | 0.1522 | 0.1477 | 0.1429 |
| 4    |            |        | 0.0625 | 0.0900 | 0.1028 | 0.1085 | 0.1106 | 0.1106 | 0.1096 |
| 5    |            |        |        | 0.0400 | 0.0611 | 0.0728 | 0.0793 | 0.0828 | 0.0846 |
| 6    |            |        |        |        | 0.0278 | 0.0442 | 0.0543 | 0.0606 | 0.0646 |
| 7    |            |        |        |        |        | 0.0204 | 0.0334 | 0.0421 | 0.0479 |
| 8    |            |        |        |        |        |        | 0.0156 | 0.0262 | 0.0336 |
| 9    |            |        |        |        |        |        |        | 0.0123 | 0.0211 |
| 10   |            |        |        |        |        |        |        |        | 0.0100 |

Abb. 13: ROC weights <sup>110</sup>

<sup>110</sup>Taken from: Roberts/Goodwin 2002

| Rank | Attributes |        |        |        |        |        |        |        |        |
|------|------------|--------|--------|--------|--------|--------|--------|--------|--------|
|      | 2          | 3      | 4      | 5      | 6      | 7      | 8      | 9      | 10     |
| 1    | 0.6667     | 0.5455 | 0.4800 | 0.4379 | 0.4082 | 0.3857 | 0.3679 | 0.3535 | 0.3414 |
| 2    | 0.3333     | 0.2727 | 0.2400 | 0.2190 | 0.2041 | 0.1928 | 0.1840 | 0.1767 | 0.1707 |
| 3    |            | 0.1818 | 0.1600 | 0.1460 | 0.1361 | 0.1286 | 0.1226 | 0.1178 | 0.1138 |
| 4    |            |        | 0.1200 | 0.1095 | 0.1020 | 0.0964 | 0.0920 | 0.0884 | 0.0854 |
| 5    |            |        |        | 0.0876 | 0.0816 | 0.0771 | 0.0736 | 0.0707 | 0.0682 |
| 6    |            |        |        |        | 0.0680 | 0.0643 | 0.0613 | 0.0589 | 0.0569 |
| 7    |            |        |        |        |        | 0.0551 | 0.0525 | 0.0505 | 0.0488 |
| 8    |            |        |        |        |        |        | 0.0460 | 0.0442 | 0.0427 |
| 9    |            |        |        |        |        |        |        | 0.0393 | 0.0379 |
| 10   |            |        |        |        |        |        |        |        | 0.0341 |

Abb. 14: RR weights <sup>111</sup>

| Rank | Attributes |        |        |        |        |        |        |        |        |
|------|------------|--------|--------|--------|--------|--------|--------|--------|--------|
|      | 2          | 3      | 4      | 5      | 6      | 7      | 8      | 9      | 10     |
| 1    | 0.6667     | 0.5000 | 0.4000 | 0.3333 | 0.2857 | 0.2500 | 0.2222 | 0.2000 | 0.1818 |
| 2    | 0.3333     | 0.3333 | 0.3000 | 0.2667 | 0.2381 | 0.2143 | 0.1944 | 0.1778 | 0.1636 |
| 3    |            | 0.1667 | 0.2000 | 0.2000 | 0.1905 | 0.1786 | 0.1667 | 0.1556 | 0.1455 |
| 4    |            |        | 0.1000 | 0.1333 | 0.1429 | 0.1429 | 0.1389 | 0.1333 | 0.1273 |
| 5    |            |        |        | 0.0667 | 0.0952 | 0.1071 | 0.1111 | 0.1111 | 0.1091 |
| 6    |            |        |        |        | 0.0476 | 0.0714 | 0.0833 | 0.0889 | 0.0909 |
| 7    |            |        |        |        |        | 0.0357 | 0.0556 | 0.0667 | 0.0727 |
| 8    |            |        |        |        |        |        | 0.0278 | 0.0444 | 0.0545 |
| 9    |            |        |        |        |        |        |        | 0.0222 | 0.0364 |
| 10   |            |        |        |        |        |        |        |        | 0.0182 |

Abb. 15: RS weights <sup>112</sup>

<sup>111</sup>Taken from: Roberts/Goodwin 2002

## Appendix 1/7: Pipeline Communication Swim Lane Diagram

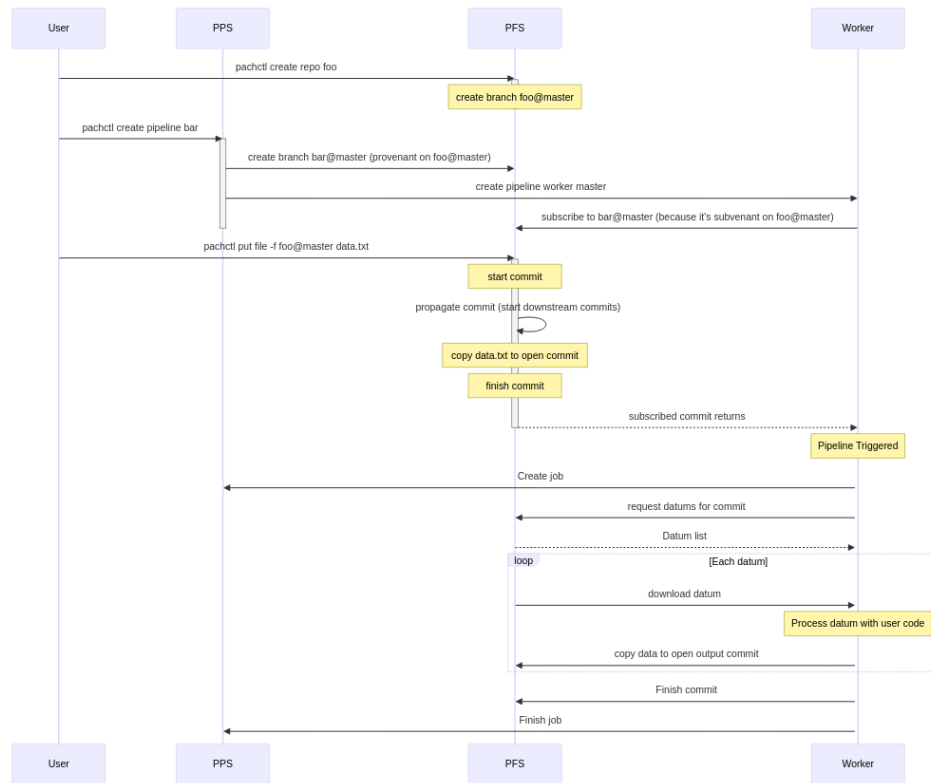


Abb. 16: Swim lane Diagram of the communication between the user and Pachyderm<sup>113</sup>

<sup>112</sup>Taken from: Roberts/Goodwin 2002

<sup>113</sup>Taken from: *Intro to Pipelines* 2023



## Appendix 2: Diagrams

### Appendix 2/1: Jenkins Pipeline Diagram

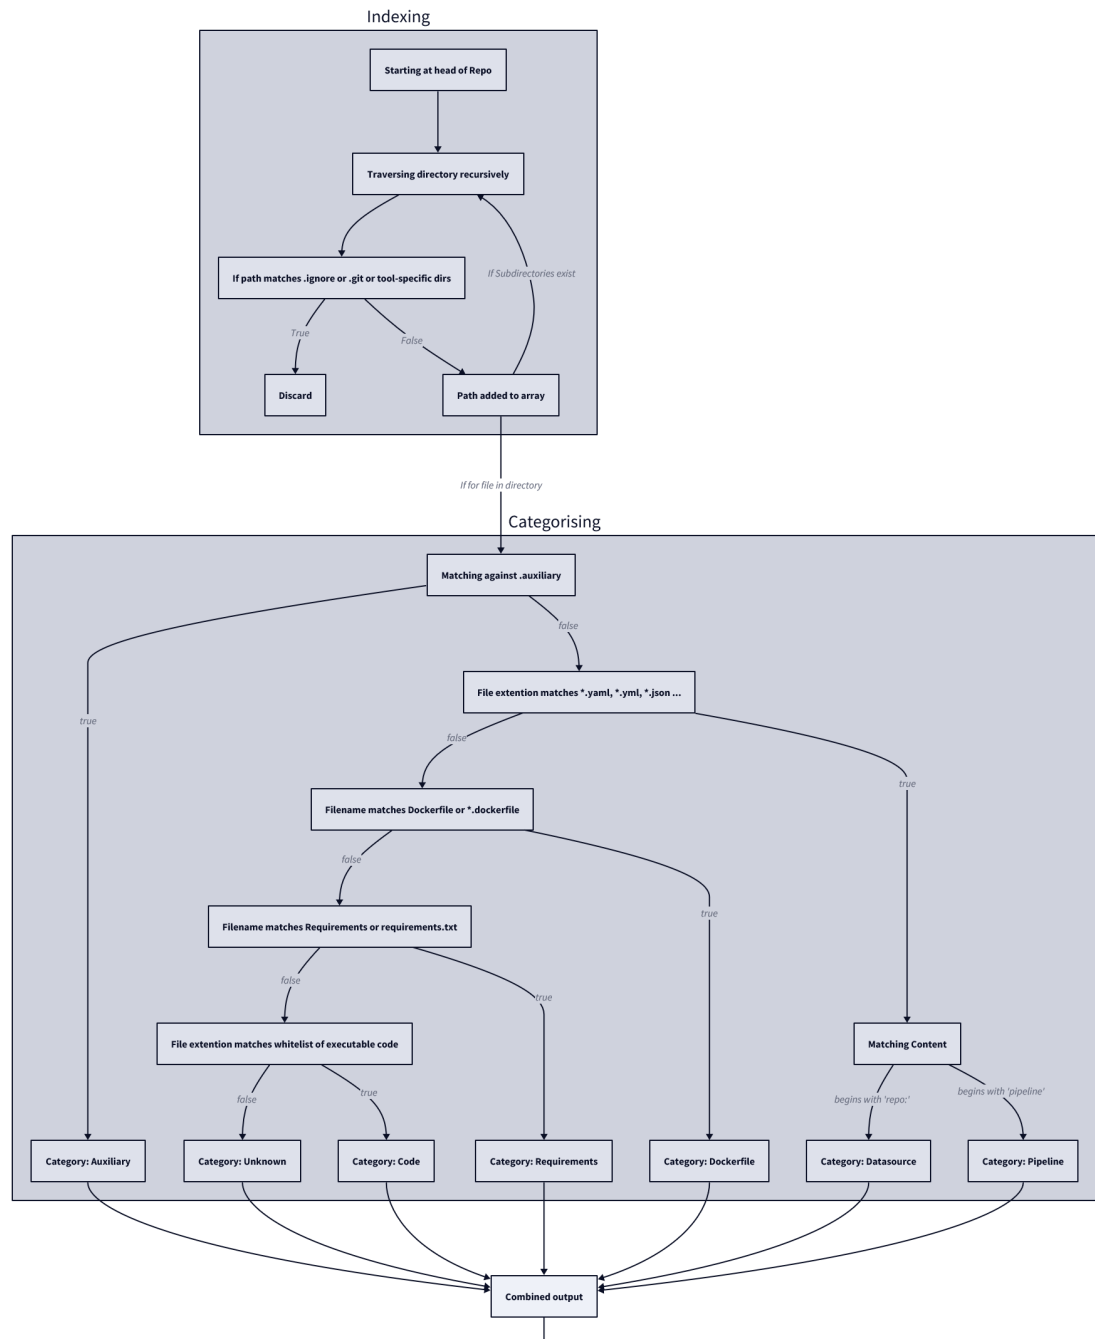


Abb. 17: Flowchart of the Jenkins Pipeline

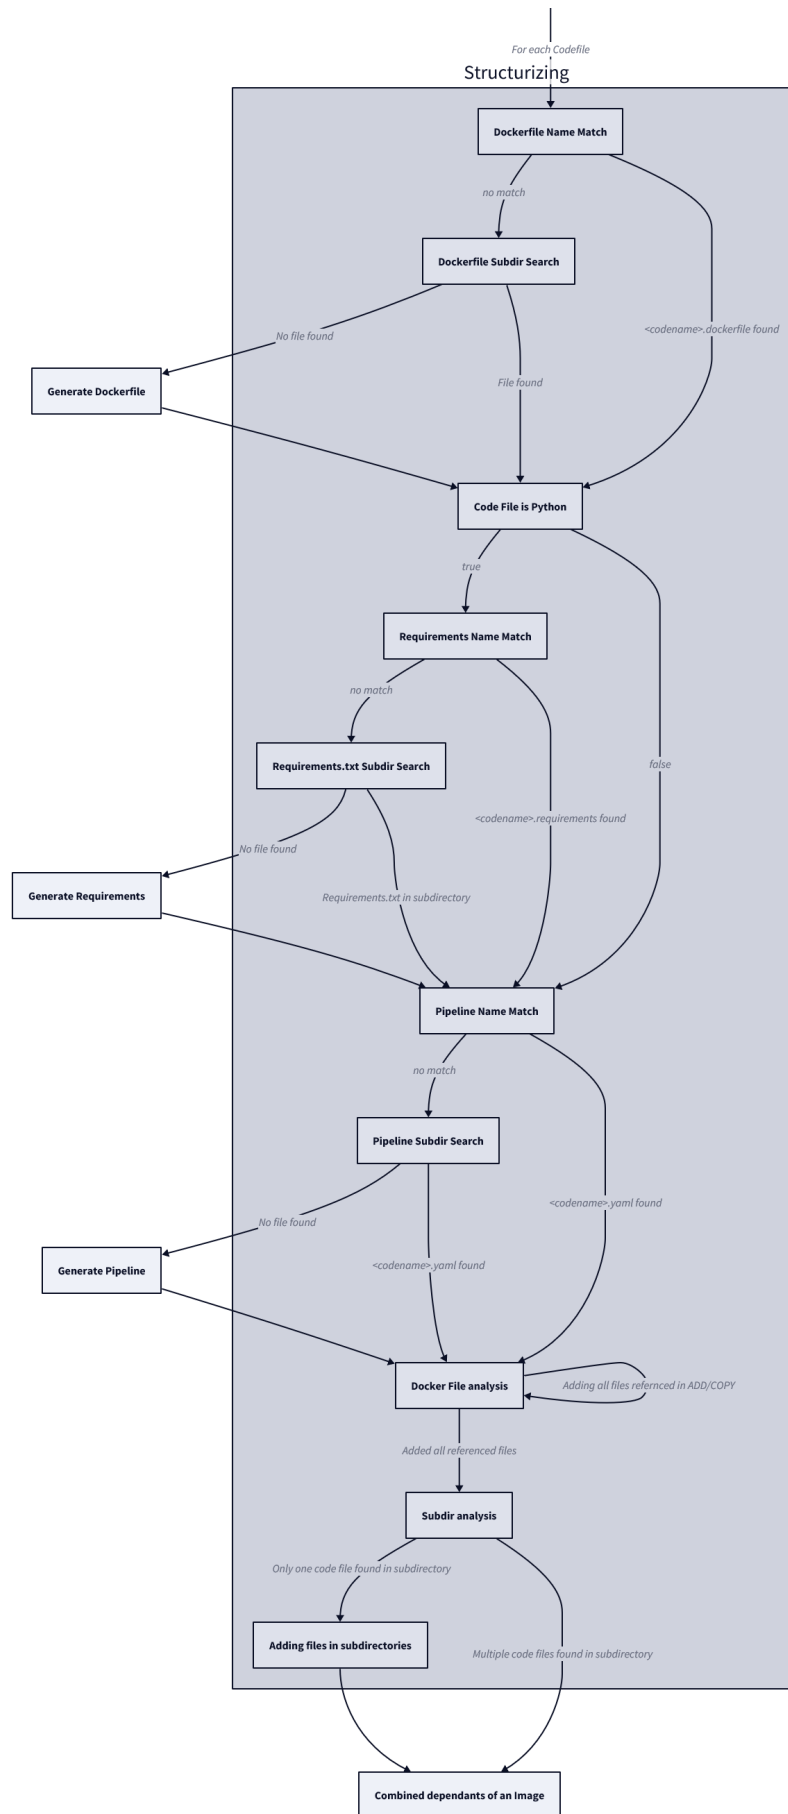


Abb. 18: Continued flowchart of the Jenkins Pipeline

## Appendix 3: Minikube installation instructions

```
1 # Pachyderm
2
3 ## Installation
4
5 These instructions are based upon the excellent guide by
6     ↪ [Pachyderm](https://docs.pachyderm.com/latest/set-up/on-prem/)
7
8
9 ### Proxy
10
11 If you are in the HPE internal network, you will need to set up the proxy.
12 Simply execute the following command:
13
14 ```bash
15 export HTTP_PROXY=http://web-proxy.corp.hpecorp.net:8080
16 export HTTPS_PROXY=http://web-proxy.corp.hpecorp.net:8080
17 ```
18
19 If you want to make this permanent, add these lines to the '~/.bashrc' or
20     ↪ equivalent file.
21
22
23 ### kubectl
24
25 Simply following the instructions on the [kubernetes
26     ↪ website](https://kubernetes.io/docs/tasks/tools/install-kubectl-linux/)
27     ↪ should be sufficient.
28
29 But for the sake of completeness, here is what I did:
30
31 ```bash
32 curl -LO "https://dl.k8s.io/release/$(curl -L -s
33     ↪ https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
34 sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl
35 ```
36
37 If the proxy is giving you grief one can simply download the binary elsewhere
38     ↪ and copy it to the target machine. (not recommended)
39
40
41 ### Installing minikube
42
43 The same things apply for minikube as for kubectl.
44 The proper instructions can be found on the [minikube
45     ↪ website](https://minikube.sigs.k8s.io/docs/start/)
46
47 But here is what I did anyway:
48
49 ```bash
50 curl -LO
51     ↪ https://storage.googleapis.com/minikube/releases/latest/minikube_latest_amd64.deb
52 sudo dpkg -i minikube_latest_amd64.deb
53 ```
54
```

```
42 We can then test the installation by running:
43
44 ```bash
45 minikube start
46 kubectl cluster-info
47 ```
48
49 If you are getting an error stating that it is not able to connect to the
    ↪ cluster you might need to set the following environment variable:
50
51 ```bash
52 export
    ↪ NO_PROXY=localhost,127.0.0.1,10.96.0.0/12,192.168.59.0/24,192.168.49.0/24,192.168.
53 ```
54
55 ### Installing [helm](https://helm.sh/docs/intro/install/)
56
57 Same procedure as every year...
58
59 ```bash
60 curl https://baltocdn.com/helm/signing.asc | gpg --dearmor | sudo tee
    ↪ /usr/share/keyrings/helm.gpg > /dev/null
61 sudo apt-get install apt-transport-https --yes
62 echo "deb [arch=$(dpkg --print-architecture)
    ↪ signed-by=/usr/share/keyrings/helm.gpg]
    ↪ https://baltocdn.com/helm/stable/debian/ all main" | sudo tee
    ↪ /etc/apt/sources.list.d/helm-stable-debian.list
63 sudo apt-get update
64 sudo apt-get install helm
65 ```
66
67 ### [Persistent
    ↪ Storage](https://kubernetes.io/docs/tasks/configure-pod-container/configure-persis
68
69 We need to create a persistent volume for etcd and the postgres database.
70 Therefore we need to create a directory for each of them.
71
72 ```bash
73 mkdir -p /mnt/pachyderm/etcd
74 mkdir -p /mnt/pachyderm/postgres
75 ```
76
77 We then create the configuration files for the persistent volumes.
78
79 ```yaml
80 apiVersion: v1
81 kind: PersistentVolume
82 metadata:
83   name: etcd-pv
84 labels:
85   type: local
```

```
86 spec:
87   capacity:
88     storage: 10Gi
89   accessModes:
90     - ReadWriteOnce
91   storageClassName: manual
92   local:
93     path: /mnt/pachyderm/etcd
94
95 ---
96
97 apiVersion: v1
98 kind: PersistentVolume
99 metadata:
100   name: postgres-pv
101 labels:
102   type: local
103 spec:
104   capacity:
105     storage: 10Gi
106   accessModes:
107     - ReadWriteOnce
108   storageClassName: manual
109   local:
110     path: /mnt/pachyderm/postgres
111 ```
112
113 And then the corresponding persistent volume claims.
114
115 ```yaml
116 apiVersion: v1
117 kind: PersistentVolumeClaim
118 metadata:
119   name: etcd-pvc
120 spec:
121   storageClassName: manual
122   accessModes:
123     - ReadWriteOnce
124   resources:
125     requests:
126       storage: 10Gi
127
128 ---
129
130 apiVersion: v1
131 kind: PersistentVolumeClaim
132 metadata:
133   name: postgres-pvc
134 spec:
135   storageClassName: manual
136   accessModes:
```

```

137         - ReadWriteOnce
138     resources:
139         requests:
140         storage: 10Gi
141     '''
142
143 Then we add the storage class to the cluster.
144
145 '''bash
146 kubectl apply -f filename.yaml
147 '''
148
149 We then take note of the storage class name because we will add it to the helm
    ↪ values file later. \
150 In this case it is 'manual'.
151
152 ### Installing [MinIO](https://min.io/docs/minio/linux/index.html)
153
154 We now install an S3 compatible storage system. Which one does not really
    ↪ matter, but I chose MinIO because it is easy to install and configure.
155
156 '''bash
157 wget
    ↪ https://dl.min.io/server/minio/release/linux-amd64/archive/minio_20230619195250.0.
    ↪ -0 minio.deb
158 sudo dpkg -i minio.deb
159
160 mkdir -p /mnt/pachyderm/minio
161
162 # to manually start the server
163 minio server /mnt/pachyderm/minio --console-address :9001
164 '''
165
166 The standard password is 'minioadmin:minioadmin'
167
168 Then you can access the web interface at 'http://localhost:9001' where you
    ↪ should login, change the password and create a bucket. \
169 The access credentials for the bucket will be added to the helm values file
    ↪ later, so take note of them.
170
171 ### Installing [Pachyderm](https://docs.pachyderm.com/latest/set-up/on-prem/)
172
173 First we need to add the Pachyderm helm repository:
174
175 '''bash
176 helm repo add pachyderm https://helm.pachyderm.com
177 helm repo update
178 '''
179
180 We then get the values file from the repository and edit it to our liking.\
181 My setup is based on the version '2.6.4-1', so it might be different for future

```

```

    ↪ versions.
182
183 ```bash
184 wget
    ↪ https://raw.githubusercontent.com/pachyderm/pachyderm/2.6.x/etc/helm/pachyderm/val
185 ```
186
187 ##### MinIO
188
189 First we change the deploy target at line 'L7'
190
191 ```yaml
192 # Deploy Target configures the storage backend to use and cloud provider
193 # settings (storage classes, etc). It must be one of GOOGLE, AMAZON,
194 # MINIO, MICROSOFT, CUSTOM or LOCAL.
195 deployTarget: "MINIO"
196 ...
197 ```
198
199 This does not need to be set when using something else but with MinIO we also
    ↪ have to set 'L544' to "MINIO"
200
201 ```yaml
202 ...
203 storage:
204     # backend configures the storage backend to use. It must be one
205     # of GOOGLE, AMAZON, MINIO, MICROSOFT or LOCAL. This is set automatically
206     # if deployTarget is GOOGLE, AMAZON, MICROSOFT, or LOCAL
207     backend: "MINIO"
208     ...
209 ```
210
211 A little further down ('L635') we find the MinIO configuration. We need to set
    ↪ the endpoint, access key and secret key.
212
213 This point was a little tricky as I had MinIO installed on the same machine as
    ↪ Pachyderm, but it would take no other value than the outward facing IP
    ↪ address of the machine.
214
215 ```yaml
216 ...
217     minio:
218         # minio bucket name
219         bucket: "<bucket name>"
220         # the minio endpoint. Should only be the hostname:port, no http/https.
221         endpoint: "10.X.X.X:9000"
222         # the username/id with readwrite access to the bucket.
223         id: "<id>"
224         # the secret/password of the user with readwrite access to the bucket.
225         secret: "<secret>"
226         # enable https for minio with "true" defaults to "false"

```

```
227     secure: "false"
228     # Enable S3v2 support by setting signature to "1". This feature is being
229     ↪ deprecated
230     signature: ""
231     ...
232
233 ##### Storage classes
234
235 Now we add the storage classes we created earlier to the Postgres at 'L784'
236
237 ```yaml
238 ...
239     # AWS: https://docs.aws.amazon.com/eks/latest/userguide/storage-classes.html
240     # GCP: https://cloud.google.com/compute/docs/disks/performance#disk_types
241     # Azure:
242     ↪ https://docs.microsoft.com/en-us/azure/aks/concepts-storage#storage-classes
243     storageClass: manual
244     # storageSize specifies the size of the volume to use for postgresql
245     # Recommended Minimum Disk size for Microsoft/Azure: 256Gi - 1,100 IOPS
246     ↪ https://azure.microsoft.com/en-us/pricing/details/managed-disks/
247     ...
248
249 and for the etcd at around 'L144'
250
251 ```yaml
252 ...
253     # GCP: https://cloud.google.com/compute/docs/disks/performance#disk_types
254     # Azure:
255     ↪ https://docs.microsoft.com/en-us/azure/aks/concepts-storage#storage-classes
256     #storageClass: manual
257     storageClassName: manual
258
259     # storageSize specifies the size of the volume to use for etcd.
260     # Recommended Minimum Disk size for Microsoft/Azure: 256Gi - 1,100 IOPS
261     ↪ https://azure.microsoft.com/en-us/pricing/details/managed-disks/
262     ...
263
264 ##### SSL Certificates
265
266 My setup refuses to work without SSL certificates, so I had to generate some.
267
268 ```bash
269 openssl genrsa -out <CertName>.key 2048
270 openssl req -new -x509 -sha256 -key <CertName>.key -out <CertName>.crt
271
272 kubectl create secret tls <SecretName> --cert=<CertName>.crt
```



```

    ↪ --key=<CertName>.key
273 ```
274
275 We then edit the 'values.yaml' file at around 'L683' to use the certificates.
276
277 ```yaml
278 ...
279   tls:
280     enabled: true
281     secretName: "<SecretName>"
282     newSecret:
283       create: false
284     ...
285 ```
286
287 ### CLI
288
289 To directly interact with the cluster we need to install the Pachyderm CLI.
290
291 ```bash
292 curl -o /tmp/pachctl.deb -L
    ↪ https://github.com/pachyderm/pachyderm/releases/download/v2.6.5/pachctl_2.6.5_amd64.deb
    ↪ && sudo dpkg -i /tmp/pachctl.deb
293 ```
294
295 ### Deploy
296
297 Now that the values file is ready we can install Pachyderm.
298
299 ```bash
300 helm install pachyderm pachyderm/pachyderm \
301   -f ./values.yaml pachyderm/pachyderm \
302   --set postgresql.volumePermissions.enabled=true \
303   --set deployTarget=LOCAL \
304   --set proxy.enabled=true \
305   --set proxy.service.type=NodePort \
306   --set proxy.host=localhost \
307   --set proxy.service.httpPort=8080
308
309 ```
310
311 Now you might want to connect to the dashboard. This can be done by
    ↪ port-forwarding the service.
312
313 ```bash
314 pachctl port-forward
315 ```
316
317 :tada: Now we should be able to access the dashboard at 'http://localhost:4000'
    ↪ :tada:

```

## Appendix 4: CI/CD installation instructions

```
1 # Setting up the GitOps CI/CD
2
3 This part is concerned with setting up both the Version Control System Forgejo
  ↪ and the CI/CD system Jenkins. \
4 While Forgejo is a fork of Gitea, it is still sparsely documented and thus we
  ↪ will not setup the runner system of Forgejo, but instead use Jenkins for
  ↪ CI/CD.
5
6
7 ## Namespaces
8
9 As always we create namespaces to keep things clean:
10
11 ```bash
12 kubectl create namespace forgejo
13 kubectl create namespace jenkins
14 ```
15
16 ## Persistent Volumes
17
18 We need to create a location for the persistent volume:
19
20 ```bash
21 mkdir -p /mnt/forgejo/postgres
22 mkdir -p /mnt/forgejo/zero
23 mkdir -p /mnt/jenkins
24 sudo chown -R eckerth:users /mnt/forgejo/ /mnt/jenkins
25 ```
26
27 We then create the persistent volumes for Forgejo ....:
28
29 ```yaml
30 apiVersion: storage.k8s.io/v1
31 kind: StorageClass
32 metadata:
33   name: forgejo
34 provisioner: kubernetes.io/no-provisioner
35 volumeBindingMode: WaitForFirstConsumer
36 ---
37 apiVersion: v1
38 kind: PersistentVolume
39 metadata:
40   name: forgejo-postgres
41   labels:
42     type: local
43 spec:
44   capacity:
45     storage: 10Gi
46   accessModes:
```

```
47     - ReadWriteOnce
48     storageClassName: forgejo
49     local:
50       path: /mnt/forgejo/postgres
51     nodeAffinity:
52       required:
53         nodeSelectorTerms:
54         - matchExpressions:
55           - key: kubernetes.io/hostname
56             operator: In
57             values:
58           - heydar20.labs.hpecorp.net
59 ---
60 apiVersion: v1
61 kind: PersistentVolume
62 metadata:
63   name: forgejo-0
64   labels:
65     type: local
66 spec:
67   capacity:
68     storage: 10Gi
69   accessModes:
70     - ReadWriteOnce
71   storageClassName: forgejo
72   local:
73     path: /mnt/forgejo/zero
74   nodeAffinity:
75     required:
76       nodeSelectorTerms:
77       - matchExpressions:
78         - key: kubernetes.io/hostname
79           operator: In
80           values:
81         - heydar20.labs.hpecorp.net
82
83 '''
84
85 ... and for Jenkins:
86
87 '''yaml
88 apiVersion: storage.k8s.io/v1
89 kind: StorageClass
90 metadata:
91   name: jenkins
92 provisioner: kubernetes.io/no-provisioner
93 volumeBindingMode: WaitForFirstConsumer
94
95 ---
96 apiVersion: v1
97 kind: PersistentVolume
```

```

98 metadata:
99   name: jenkins
100   labels:
101     type: local
102 spec:
103   capacity:
104     storage: 10Gi
105   accessModes:
106     - ReadWriteOnce
107   storageClassName: jenkins
108   local:
109     path: /mnt/jenkins
110   nodeAffinity:
111     required:
112       nodeSelectorTerms:
113         - matchExpressions:
114             - key: kubernetes.io/hostname
115               operator: In
116               values:
117                 - heydar20.labs.hpecorp.net
118   ```
119
120   ```bash
121   kubectl -n forgejo apply -f ./forgejo/volumes.yaml
122   kubectl -n jenkins apply -f ./jenkins/volumes.yaml
123   ```
124
125   ## Installation
126
127   After these are applied we can simply install the helm chart:
128
129   ```bash
130   helm repo add jenkins https://charts.jenkins.io
131   helm repo update
132   helm install -n jenkins jenkins jenkins/jenkins -f ./jenkins/values.yaml
133   helm install -n forgejo forgejo oci://codeberg.org/forgejo-contrib/forgejo -f
134     ↪ ./forgejo/values.yaml
135   ```
136
137   ## Configuring
138
139   In order to connect both Jenkins and Forgejo we will have to adjust some
140     ↪ configurations.
141
142   1. As we want Jenkins to be able to be able spawn pods on the cluster, we will
143     ↪ need to give it the needed permissions.
144   For this one can use the service_account.yaml file in the jenkins folder.
145
146   ```bash
147   kubectl -n jenkins apply -f ./jenkins/service_account.yaml
148   ```

```

```
146
147 2. We add the following config map to Forgejo in order to allow it to send
    ↪ webhooks to out jenkins host.
148
149 ```yaml
150 additionalConfigSources:
151     - configMap:
152         name: gitea-app-ini
153 ```
154
155 ```yaml
156 apiVersion: v1
157 kind: ConfigMap
158 metadata:
159     name: gitea-app-ini
160 data:
161     webhook: |
162         ALLOWED_HOST_LIST=<jenkins server>
163 ```
164
165 ```bash
166 kubectl -n forgejo apply -f ./forgejo/configmap.yaml
167 ```
168
169 2. We then have to go into the Forgejo admin pannel and enable the system wide
    ↪ Webhooks,
```

## List of literature

- 6sense (2023)**: Jenkins - Market Share, Competitor Insights in Continuous Integration And Delivery. 6sense. <https://www.6sense.com/tech/continuous-integration/jenkins-market-share> (retrieval: 10/31/2023).
- Ansible (2023)**: *Ansible*. Ansible. <https://github.com/ansible/ansible> (retrieval: 10/22/2023).
- Apache (2023)**: Apache Airflow License. <https://airflow.apache.org/docs/apache-airflow/stable/license.html> (retrieval: 10/21/2023).
- Argo (2023a)**: *Argoproj/Argo-Workflows*. Argo Project. <https://github.com/argoproj/argo-workflows> (retrieval: 10/07/2023).
- **(2023b)**: Argo-Cd/LICENSE at Master · Argo-Cd. <https://github.com/argoproj/argo-cd/blob/master/LICENSE> (retrieval: 10/21/2023).
- Barfod, M. B./Leleur, S. (2014)**: Multi-Criteria Decision Analysis for Use in Transport Decision Making. DTU Lyngby: DTU Transport.
- Baur, D./Seybold, D./Griesinger, F./Tsitsipas, A./Hauser, C. B./Domaschka, J. (2015)**: Cloud Orchestration Features: Are Tools Fit for Purpose? In: *2015 IEEE/ACM 8th International Conference on Utility and Cloud Computing (UCC)*. 2015 IEEE/ACM 8th International Conference on Utility and Cloud Computing (UCC), pp. 95–101. 10.1109/UCC.2015.25. <https://ieeexplore.ieee.org/abstract/document/7431400> (retrieval: 11/05/2023).
- Bears-R-Us (2023a)**: *Arkouda Github Repository*. Bears-R-Us. <https://github.com/Bears-R-Us/arkouda> (retrieval: 10/23/2023).
- **(2023b)**: Arkouda-Contrib/Arkouda-Docker at Main · Bears-R-Us/Arkouda-Contrib. GitHub. <https://github.com/Bears-R-Us/arkouda-contrib/tree/main/arkouda-docker> (retrieval: 10/23/2023).
- **(2023c)**: Bears-R-Us/Arkouda-Contrib/Arkouda-Helm-Charts. GitHub. <https://github.com/Bears-R-Us/arkouda-contrib/tree/main/arkouda-helm-charts/arkouda-udp-server> (retrieval: 10/25/2023).
- Bigelow, S. J. (2023)**: Explore Network Plugins for Kubernetes: CNI Explained | TechTarget. IT Operations. <https://www.techtarget.com/searchitoperations/tip/Explore-network-plugins-for-Kubernetes-CNI-explained> (retrieval: 10/19/2023).
- Boehm, B. W. (1988)**: A Spiral Model of Software Development and Enhancement. In: *Computer* 21.5, pp. 61–72. ISSN: 0018-9162. 10.1109/2.59. <http://ieeexplore.ieee.org/document/59/> (retrieval: 10/18/2023).
- Brown, R. (2004)**: Engineering a Beowulf-Style Compute Cluster. In.
- Byrne, J./Kuno, H./Ghosh, C./Shome, P. (2023)**: Coupling Chapel-Powered HPC Workflows for Python. In.
- Bzeznik, B./Henriot, O./Reis, V./Richard, O./Tavard, L. (2017)**: Nix as HPC Package Management System. In: *Proceedings of the Fourth International Workshop on HPC User Support Tools*. HUST’17. New York, NY, USA: Association for Computing Machinery, pp. 1–6. ISBN: 978-1-4503-5130-0. 10.1145/3152493.3152556. <https://dl.acm.org/doi/10.1145/3152493.3152556> (retrieval: 10/03/2023).
- Camacho-Rodríguez, J./Chauhan, A./Gates, A./Koifman, E./O’Malley, O./Garg, V./Haindrich, Z./Shelukhin, S./Jayachandran, P./Seth, S./Jaiswal, D./Bouguerra,**

- S./Bangarwa, N./Hariappan, S./Agarwal, A./Dere, J./Dai, D./Nair, T./Dembla, N./Vijayaraghavan, G./Hagleitner, G. (2019):** Apache Hive: From MapReduce to Enterprise-grade Big Data Warehousing. In: *Proceedings of the 2019 International Conference on Management of Data*. SIGMOD '19. New York, NY, USA: Association for Computing Machinery, pp. 1773–1786. ISBN: 978-1-4503-5643-5. 10.1145/3299869.3314045. <https://dl.acm.org/doi/10.1145/3299869.3314045> (retrieval: 11/07/2023).
- Canon, R. S./Younge, A. (2019):** A Case for Portability and Reproducibility of HPC Containers. In: *2019 IEEE/ACM International Workshop on Containers and New Orchestration Paradigms for Isolated Environments in HPC (CANOPIE-HPC)*. 2019 IEEE/ACM International Workshop on Containers and New Orchestration Paradigms for Isolated Environments in HPC (CANOPIE-HPC). Denver, CO, USA: IEEE, pp. 49–54. ISBN: 978-1-72816-028-3. 10.1109/CANOPIE-HPC49598.2019.00012. <https://ieeexplore.ieee.org/document/8950982/> (retrieval: 10/04/2023).
- CephIo (2023):** Ceph.Io — Home. <https://ceph.io/en/> (retrieval: 10/22/2023).
- Chapman, B./Curtis, T./Pophale, S./Poole, S./Kuehn, J./Koelbel, C./Smith, L. (2010):** Introducing OpenSHMEM: SHMEM for the PGAS Community. In: *Proceedings of the Fourth Conference on Partitioned Global Address Space Programming Model*. PGAS '10: Fourth Conference on Partitioned Global Address Space Programming Model. New York New York USA: ACM, pp. 1–3. ISBN: 978-1-4503-0461-0. 10.1145/2020373.2020375. <https://dl.acm.org/doi/10.1145/2020373.2020375> (retrieval: 11/07/2023).
- Chou, C. C./Chen, Y./Milojicic, D./Reddy, N./Gratz, P. (2019):** Optimizing Post-Copy Live Migration with System-Level Checkpoint Using Fabric-Attached Memory. In: *2019 IEEE/ACM Workshop on Memory Centric High Performance Computing (MCHPC)*. 2019 IEEE/ACM Workshop on Memory Centric High Performance Computing (MCHPC). Denver, CO, USA: IEEE, pp. 16–24. ISBN: 978-1-72816-007-8. 10.1109/MCHPC49590.2019.00010. <https://ieeexplore.ieee.org/document/8946189/> (retrieval: 10/29/2023).
- Codeberg (2023):** Codeberg.Org. Codeberg.org. <https://codeberg.org/> (retrieval: 10/31/2023).
- Common-Workflow-Language (2023):** Cwl-Utils/LICENSE. <https://github.com/common-workflow-language/cwl-utils/blob/main/LICENSE> (retrieval: 10/21/2023).
- Docker (2021):** Docker FAQs | Docker. <https://www.docker.com/pricing/faq/> (retrieval: 10/18/2023).
- **(2023):** What Is a Container? <https://www.docker.com/resources/what-container/> (retrieval: 11/05/2023).
- Docker Terms of Service | Docker (2022).** <https://www.docker.com/legal/docker-terms-service/> (retrieval: 10/18/2023).
- Dubois, P./Epperly, T./Kumfert, G. (2003):** Why Johnny Can't Build [Portable Scientific Software]. In: *Computing in Science & Engineering* 5.5, pp. 83–88. ISSN: 1558-366X. 10.1109/MCISE.2003.1225867. <https://ieeexplore.ieee.org/abstract/document/1225867> (retrieval: 10/03/2023).

- Eckerth, J. (2023a):** Installation Instructions Minikube. <https://github.com/Jo-Eck/ProjectPaper-2/tree/69a05a755facc8d387e031ff5991c20d46dbe4b6/project/Pachyderm> (retrieval: 10/18/2023).
- **(2023b):** Project Repo Ansible Script. GitHub. <https://github.com/Jo-Eck/ProjectPaper-2> (retrieval: 11/17/2023).
  - **(2023c):** Project Repo Docker Registry. GitHub. [https://github.com/Jo-Eck/ProjectPaper-2/tree/main/Project/Kubernetes\\_Setup/registry](https://github.com/Jo-Eck/ProjectPaper-2/tree/main/Project/Kubernetes_Setup/registry) (retrieval: 11/17/2023).
  - **(2023d):** Project Repo Flannel Config. GitHub. <https://github.com/Jo-Eck/ProjectPaper-2> (retrieval: 11/17/2023).
  - **(2023e):** Project Repo Forgejo Configuration. GitHub. [https://github.com/Jo-Eck/ProjectPaper-2/tree/main/Project/Kubernetes\\_Setup/ci\\_cd/forgejo](https://github.com/Jo-Eck/ProjectPaper-2/tree/main/Project/Kubernetes_Setup/ci_cd/forgejo) (retrieval: 11/17/2023).
  - **(2023f):** Project Repo Heydar Setup. GitHub. [https://github.com/Jo-Eck/ProjectPaper-2/tree/main/Project/Kubernetes\\_Setup/kluster](https://github.com/Jo-Eck/ProjectPaper-2/tree/main/Project/Kubernetes_Setup/kluster) (retrieval: 11/17/2023).
  - **(2023g):** Project Repo Main Readme. GitHub. <https://github.com/Jo-Eck/ProjectPaper-2/blob/main/Project/README.md> (retrieval: 11/17/2023).
- Egwutuoha, I. P./Levy, D./Selic, B./Chen, S. (2013):** A Survey of Fault Tolerance Mechanisms and Checkpoint/Restart Implementations for High Performance Computing Systems. In: *J Supercomput* 65.3, pp. 1302–1326. ISSN: 1573-0484. 10.1007/s11227-013-0884-0. <https://doi.org/10.1007/s11227-013-0884-0> (retrieval: 10/03/2023).
- Enterprise, H. P. (2023):** Hewlett Packard Enterprise Acquires Pachyderm to Expand AI-at-scale Capabilities with Reproducible AI. <https://www.hpe.com/us/en/newsroom/press-release/2023/01/hewlett-packard-enterprise-acquires-pachyderm-to-expand-ai-at-scale-capabilities-with-reproducible-ai.html> (retrieval: 11/07/2023).
- Flannel (2023a):** *Flannel*. flannel-io. <https://github.com/flannel-io/flannel> (retrieval: 10/19/2023).
- **(2023b):** Flannel Install Config. <https://github.com/flannel-io/flannel/blob/master/Documentation/kube-flannel.yml> (retrieval: 10/19/2023).
- Forbes advisor (2023):** 24 Top AI Statistics & Trends In 2023 – Forbes Advisor. <https://www.forbes.com/advisor/business/ai-statistics/> (retrieval: 10/07/2023).
- Forgejo (2023a):** Forgejo. Codeberg.org. <https://codeberg.org/forgejo/forgejo> (retrieval: 10/30/2023).
- **(2023b):** Forgejo 0.13.0 · Forgejo/Forgejo-Helm. <https://artifacthub.io/packages/helm/forgejo-helm/forgejo> (retrieval: 10/31/2023).
- Fülöp, J. (2005):** Introduction to Decision Making Methods. In: *BDEI-3 Workshop, Washington*, pp. 1–15. <https://www.academia.edu/download/43447287/decisionmakingmethods.pdf> (retrieval: 10/12/2023).
- Gamblin, T./LeGendre, M./Collette, M. R./Lee, G. L./Moody, A./de Supinski, B. R./Futral, S. (2015):** The Spack Package Manager: Bringing Order to HPC Software Chaos. In: *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*. SC '15. New York, NY, USA: Association for Computing Machinery, pp. 1–12. ISBN: 978-1-4503-3723-6. 10.1145/2807591.2807623. <https://dl.acm.org/doi/10.1145/2807591.2807623> (retrieval: 10/03/2023).



- Gitea (2023)**: *Gitea - Git with a Cup of Tea*. Gitea. <https://github.com/go-gitea/gitea> (retrieval: 10/31/2023).
- Gluster (2023)**: Gluster. <https://www.gluster.org/> (retrieval: 10/22/2023).
- Gomaa, H. (1983)**: The Impact of Rapid Prototyping on Specifying User Requirements. In: *SIGSOFT Softw. Eng. Notes* 8.2, pp. 17–27. ISSN: 0163-5948. 10.1145/1005959.1005964. <http://doi.org/10.1145/1005959.1005964> (retrieval: 11/03/2023).
- Google (2023)**: *Kaniko - Build Images In Kubernetes*. GoogleContainerTools. <https://github.com/GoogleContainerTools/kaniko> (retrieval: 11/01/2023).
- Haines, S. (2022)**: ‘Workflow Orchestration with Apache Airflow’. In: *Modern Data Engineering with Apache Spark: A Hands-On Guide for Building Mission-Critical Streaming Applications*. Springer, pp. 255–295.
- Helm (2023)**: Helm Docs Home. <https://helm.sh/docs/> (retrieval: 10/18/2023).
- Hewlett Packard Enterprise Development LP (2023)**: *Chapel-Lang/Chapel: A Productive Parallel Programming Language*. <https://github.com/chapel-lang/chapel> (retrieval: 10/23/2023).
- Hoste, K./Timmerman, J./Georges, A./De Weirdt, S. (2012)**: EasyBuild: Building Software with Ease. In: *2012 SC Companion: High Performance Computing, Networking Storage and Analysis*. 2012 SC Companion: High Performance Computing, Networking Storage and Analysis, pp. 572–582. 10.1109/SC.Companion.2012.81. <https://ieeexplore.ieee.org/abstract/document/6495863> (retrieval: 10/03/2023).
- Inc, M. (2023)**: MinIO | MinIO for Kubernetes. MinIO. <https://min.io> (retrieval: 10/18/2023).
- Intro to Pipelines (2023)**. <https://docs.pachyderm.com/latest/learn/intro-pipelines/> (retrieval: 10/23/2023).
- Jenkins (2023a)**: Jenkins 4.8.2 · Jenkins/Jenkinsci. <https://artifacthub.io/packages/helm/jenkinsci/jenkins> (retrieval: 10/31/2023).
- **(2023b)**: Jenkins/Jenkins - Docker Image | Docker Hub. <https://hub.docker.com/r/jenkins/jenkins> (retrieval: 11/01/2023).
- Jin, H./Chen, Y./Zhu, H./Sun, X.-H. (2010)**: Optimizing HPC Fault-Tolerant Environment: An Analytical Approach. In: *2010 39th International Conference on Parallel Processing*. 2010 39th International Conference on Parallel Processing (ICPP). San Diego, CA, USA: IEEE, pp. 525–534. ISBN: 978-1-4244-7913-9. 10.1109/ICPP.2010.80. <http://ieeexplore.ieee.org/document/5599253/> (retrieval: 11/03/2023).
- Keeton, K./Singhal, S./Raymond, M. (2019)**: The OpenFAM API: A Programming Model for Disaggregated Persistent Memory. In: *OpenSHMEM and Related Technologies. OpenSHMEM in the Era of Extreme Heterogeneity*. Ed. by Swaroop Pophale/Neena Imam/Ferrol Aderholdt/Manjunath Gorentla Venkata. Lecture Notes in Computer Science. Cham: Springer International Publishing, pp. 70–89. ISBN: 978-3-030-04918-8. 10.1007/978-3-030-04918-8\_5.
- Knative (2023a)**: Knative Docs/LICENSE at Main · Knative/Docs. <https://github.com/knative/docs/blob/main/LICENSE> (retrieval: 10/21/2023).
- **(2023b)**: Knative Homepage. <https://knative.dev/docs/> (retrieval: 10/07/2023).
- Kraushaar, J. M./Shirland, L. E. (1985)**: A Prototyping Method for Applications Development by End Users and Information Systems Specialists. In: *MIS Quarterly* 9.3, pp. 189–

197. ISSN: 0276-7783. 10.2307/248948. JSTOR: 248948. <https://www.jstor.org/stable/248948> (retrieval: 11/03/2023).
- Kubeflow (2023). Kubeflow. <https://www.kubeflow.org/> (retrieval: 10/07/2023).
- Kubeflow (2023)**: Kubeflow/LICENSE at Master · Kubeflow/Kubeflow. <https://github.com/kubeflow/kubeflow/blob/master/LICENSE> (retrieval: 10/21/2023).
- Kubernetes (2023a)**: Cluster Networking. Kubernetes. <https://kubernetes.io/docs/concepts/cluster-administration/networking/> (retrieval: 10/18/2023).
- **(2023b)**: Creating a Cluster with Kubeadm. Kubernetes. <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/> (retrieval: 10/18/2023).
- **(2023c)**: Installing Kubeadm. Kubernetes. <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/> (retrieval: 10/22/2023).
- Kubernetes CNI Plugins (2023). Kubernetes. <https://kubernetes.io/docs/concepts/cluster-administration/addons/> (retrieval: 10/19/2023).
- Kumar, P. (2020)**: How to Setup Private Docker Registry in Kubernetes (K8s). <https://www.linuxtechi.com/setup-private-docker-registry-kubernetes/> (retrieval: 10/29/2023).
- Luigi/LICENSE at Master · Spotify/Luigi (2023). <https://github.com/spotify/luigi/blob/master/LICENSE> (retrieval: 10/21/2023).
- Martin, J. (1991)**: Rapid Application Development. In collab. with Internet Archive. New York : Macmillan Pub. Co. ; Toronto : Collier Macmillan Canada ; New York : Maxwell Macmillan International. 826 pp. ISBN: 978-0-02-376775-3. <http://archive.org/details/rapidapplication00mart> (retrieval: 11/03/2023).
- Mehndiratta, H. (2023)**: Comparing Kubernetes Container Network Interface (CNI) Providers | Kubevious.Io. <https://kubevious.io/blog/post/comparing-kubernetes-container-network-interface-cni-providers> (retrieval: 10/19/2023).
- Merrill, M./Reus, W./Neumann, T. (2019)**: Arkouda: Interactive Data Exploration Backed by Chapel. In: *Proceedings of the ACM SIGPLAN 6th on Chapel Implementers and Users Workshop*. CHIUIW 2019. New York, NY, USA: Association for Computing Machinery, p. 28. ISBN: 978-1-4503-6800-1. 10.1145/3329722.3330148. <https://dl.acm.org/doi/10.1145/3329722.3330148> (retrieval: 10/04/2023).
- N/A (w.y.)**: ‘Introduction to Software Containers with Singularity.’ (RRZE). <https://hpc.fau.de/files/2020/05/2020-05-12-singularity-containers.pdf>.
- Naumann, J. D./Jenkins, A. M. (1982)**: Prototyping: The New Paradigm for Systems Development. In: *MIS Quarterly* 6.3, pp. 29–44. ISSN: 0276-7783. 10.2307/248654. JSTOR: 248654. <https://www.jstor.org/stable/248654> (retrieval: 11/03/2023).
- Openstack (2023)**: Open Source Cloud Computing Infrastructure. OpenStack. <https://www.openstack.org/> (retrieval: 11/05/2023).
- Pachyder (2023a)**: Pachyderm Docs - Local Deploy. Pachyderm Docs - Local Deploy. <https://docs.pachyderm.com/2.6.x/set-up/local-deploy/> (retrieval: 10/18/2023).
- **(2023b)**: Pachyderm Docs - On-Prem Deploy. Pachyderm Docs - On-Prem Deploy. <https://docs.pachyderm.com/latest/set-up/on-prem/> (retrieval: 10/18/2023).

- Pachyderm (2022a):** Pachyderm Homepage. Pachyderm. <https://www.pachyderm.com/> (retrieval: 10/04/2023).
- **(2022b):** Pachyderm -Pricing. <https://www.pachyderm.com/pricing/> (retrieval: 10/21/2023).
  - **(2023a):** Pachyderm Target Audience. <https://docs.pachyderm.com/latest/learn/target-audience/> (retrieval: 11/01/2023).
  - **(2023b):** Artifacthub Pachyderm 2.6.4. <https://artifacthub.io/packages/helm/pachyderm/pachyderm/2.6.4> (retrieval: 10/18/2023).
  - **(2023c):** Pachyderm License. <https://github.com/pachyderm/pachyderm/blob/master/LICENSE> (retrieval: 10/21/2023).
- Petazzoni, J. (2023):** Using Docker-in-Docker for Your CI or Testing Environment? Think Twice. <https://jpetazzo.github.io/2015/09/03/do-not-use-docker-in-docker-for-ci/> (retrieval: 11/01/2023).
- Phntom (2023):** Docker-Registry 1.10.0 · Phntom/Phntom. <https://artifacthub.io/packages/helm/phntom/docker-registry> (retrieval: 10/30/2023).
- Production-Grade Container Orchestration (2023). Kubernetes. <https://kubernetes.io/> (retrieval: 11/05/2023).
- Red Hat OpenShift Enterprise Kubernetes Container Platform (2023). <https://www.redhat.com/en/technologies/cloud-computing/openshift> (retrieval: 11/05/2023).
- Roberts, R./Goodwin, P. (2002):** Weight Approximations in Multi-Attribute Decision Models. In: *Journal of Multi-Criteria Decision Analysis* 11.6, pp. 291–303. ISSN: 1099-1360. 10.1002/mcda.320. <https://onlinelibrary.wiley.com/doi/abs/10.1002/mcda.320> (retrieval: 10/16/2023).
- Rook (2023):** Rook. <https://rook.io/> (retrieval: 10/22/2023).
- Sayers, C./Laffitte, H./Reddy, P./Ozonat, K./Sayal, M./Simitsis, A./Singhal, S./Koutrika, G./Das, M./Aji, A./Bosamiya, H./Riss, M./Wilkinson, K./Lucio, J./Cantal, A./Carvalho, C. (2015):** Cloud Application Services Platform (CLASP): User Guide, Introduction, and Operation. In: pp. 1–60.
- Sayers, Craig/Laffitte, H./Reddy, P./Ozonat, K./Sayal, M./Simitsis, A./Singhal, S./Koutrika, G./Das, M./Aji, A./Bosamiya, H. A./Riss, M./Wilkinson, K. (2015):** CCloud Application Services Platform. In.
- Shenoi, V. V./Shah, V./Joshi, S. K. (2019):** HPC Education for Domain Scientists: An Indian Experience and Perspective. In: *2019 26th International Conference on High Performance Computing, Data and Analytics Workshop (HiPCW)*. 2019 26th International Conference on High Performance Computing, Data and Analytics Workshop (HiPCW). Hyderabad, India: IEEE, pp. 64–70. ISBN: 978-1-72814-894-6. 10.1109/HiPCW.2019.00023. <https://ieeexplore.ieee.org/document/9001797/> (retrieval: 11/01/2023).
- SnapLogic (2023):** iPaaS Solution for the Enterprise. SnapLogic. <https://www.snaplogic.com/> (retrieval: 10/07/2023).
- SnapLogic – Master Subscription Agreement (2023). <https://www.snaplogic.com/master-subscription-agreement> (retrieval: 10/21/2023).
- Snir, M. (1998):** MPI—the Complete Reference: The MPI Core. MIT Press. 452 pp. ISBN: 978-0-262-69215-1. Google Books: x79puJ2YkroC.

- Spotify (2023):** *Spotify/Luigi*. Spotify. <https://github.com/spotify/luigi> (retrieval: 10/12/2023).
- StackOverflow (2023):** Stack Overflow Developer Survey 2022. Stack Overflow. <https://survey.stackoverflow.co/2022/#section-most-popular-technologies-other-tools> (retrieval: 11/05/2023).
- Triantaphyllou, E. (2000):** ‘Introduction to Multi-Criteria Decision Making’. In: *Multi-Criteria Decision Making Methods: A Comparative Study*. Ed. by Evangelos Triantaphyllou. Applied Optimization. Boston, MA: Springer US, pp. 1–4. ISBN: 978-1-4757-3157-6. 10.1007/978-1-4757-3157-6\_1. [https://doi.org/10.1007/978-1-4757-3157-6\\_1](https://doi.org/10.1007/978-1-4757-3157-6_1) (retrieval: 10/12/2023).
- Weighted Sum Model (2022).** In: *Wikipedia*. [https://en.wikipedia.org/w/index.php?title=Weighted\\_sum\\_model&oldid=1114224941](https://en.wikipedia.org/w/index.php?title=Weighted_sum_model&oldid=1114224941) (retrieval: 10/16/2023).
- Wilde, T./Hess, T. (2007):** Forschungsmethoden der Wirtschaftsinformatik. In: *WIRTSCHAFTSINFORMATIK* 49 (4). 10.1007/s11576-007-0064-z, pp. 280–287. ISSN: 0937-6429. <http://dx.doi.org/10.1007/s11576-007-0064-z>.
- Xia, W./Wen, Y./Foh, C. H./Niyato, D./Xie, H. (2015):** A Survey on Software-Defined Networking. In: *IEEE Communications Surveys & Tutorials* 17.1, pp. 27–51. ISSN: 1553-877X. 10.1109/COMST.2014.2330903. <https://ieeexplore.ieee.org/abstract/document/6834762> (retrieval: 11/05/2023).